

AlphaOthello

Due: 6/28 11:59am

Demo: 6/28 2pm

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading


Introduction

- Design and implement an AI that can play the boardgame Othello
- Read the current board and output the next move
- Design a **state value function** to evaluate the score of the board
- Determine the next move with tree search algorithm

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

Othello (Reversi)



Source: <https://www.eothello.com/>
<https://www.othelloonline.org/> (Stronger)


Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- The opposite color discs between player discs will be flipped
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- The opposite color discs between player discs will be flipped
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

Initial board




Source: <https://www.eothello.com/>

Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- The opposite color discs between player discs will be flipped
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

Gray discs are valid moves of black




Source: <https://www.eothello.com/>

Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- **The opposite color discs between player discs will be flipped**
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

Black plays (2, 3), one white disc is flipped




Source: <https://www.eothello.com/>

Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- The opposite color discs between player discs will be flipped
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

Since neither player can move, game ends




Source: <https://www.eothello.com/>

Othello Rules

- Use black and white discs and an 8 x 8 board, black plays first
- Starts with initial board (the board on the previous page)
- Player can place disc if the new disc and another same color disc are on a line (horizontal, vertical, diagonal), with one or more contiguous opposite color discs between the two discs of the player
- The opposite color discs between player discs will be flipped
- Player passes if no valid move can be performed
- The game ends when neither player can move
- The one with more discs wins

White wins with more discs



Source: <https://www.eothello.com/>

Outline

1. Introduction
2. Othello
- 3. State Value Function**
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

State Value Function

- The program should decide which move is better
- We can pick the move which leads to the board with highest score
- Thus, we need a function to evaluate the score of the board
- It is the “**state value function**”


State Value Function

- State \Rightarrow the board
- Value \Rightarrow how “good” the board is
- Function \Rightarrow given a board, output the value

Simple Example


- According to the Othello rule, we win if we have more discs than the opponent when the game ends
- Thus, we can define a simple value function:
- Value = # of player's discs – # of opponent's discs

Suppose we play as black, value = $4 - 1 = 3$



Source: <https://www.eothello.com/>


Another example, value = $40 - 24 = 16$



Source: <https://www.eothello.com/>


Use value function to pick the next move

- Suppose we have three valid moves, A, B, and C:




Use value function to pick the next move

- After evaluating the state values, we have 20, -15, and 30



Use value function to pick the next move

- We pick move C to be our next step since it leads to the highest value



Stronger Value Functions

- The value function in the previous example is too simple
- It is nearly impossible to beat the baselines by using it
- In this mini project, you must design a stronger value function to evaluate how “good” the board is

Features You Can Utilize

- Disc count
- Valid move count
- Disc position
- Game status (win, lose, skipped turn)
- Try to figure out more features by yourself

Outline

1. Introduction
2. Othello
3. State Value Function
- 4. Minimax**
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

Minimax

- In the previous example, we only look forward for one step
- However, the opponent will try its best to defeat you
- Greedy choice (next highest value) is not always the best
- We should look forward for more steps and simulate how the opponent thinks to make the best choice with least risk

Minimax

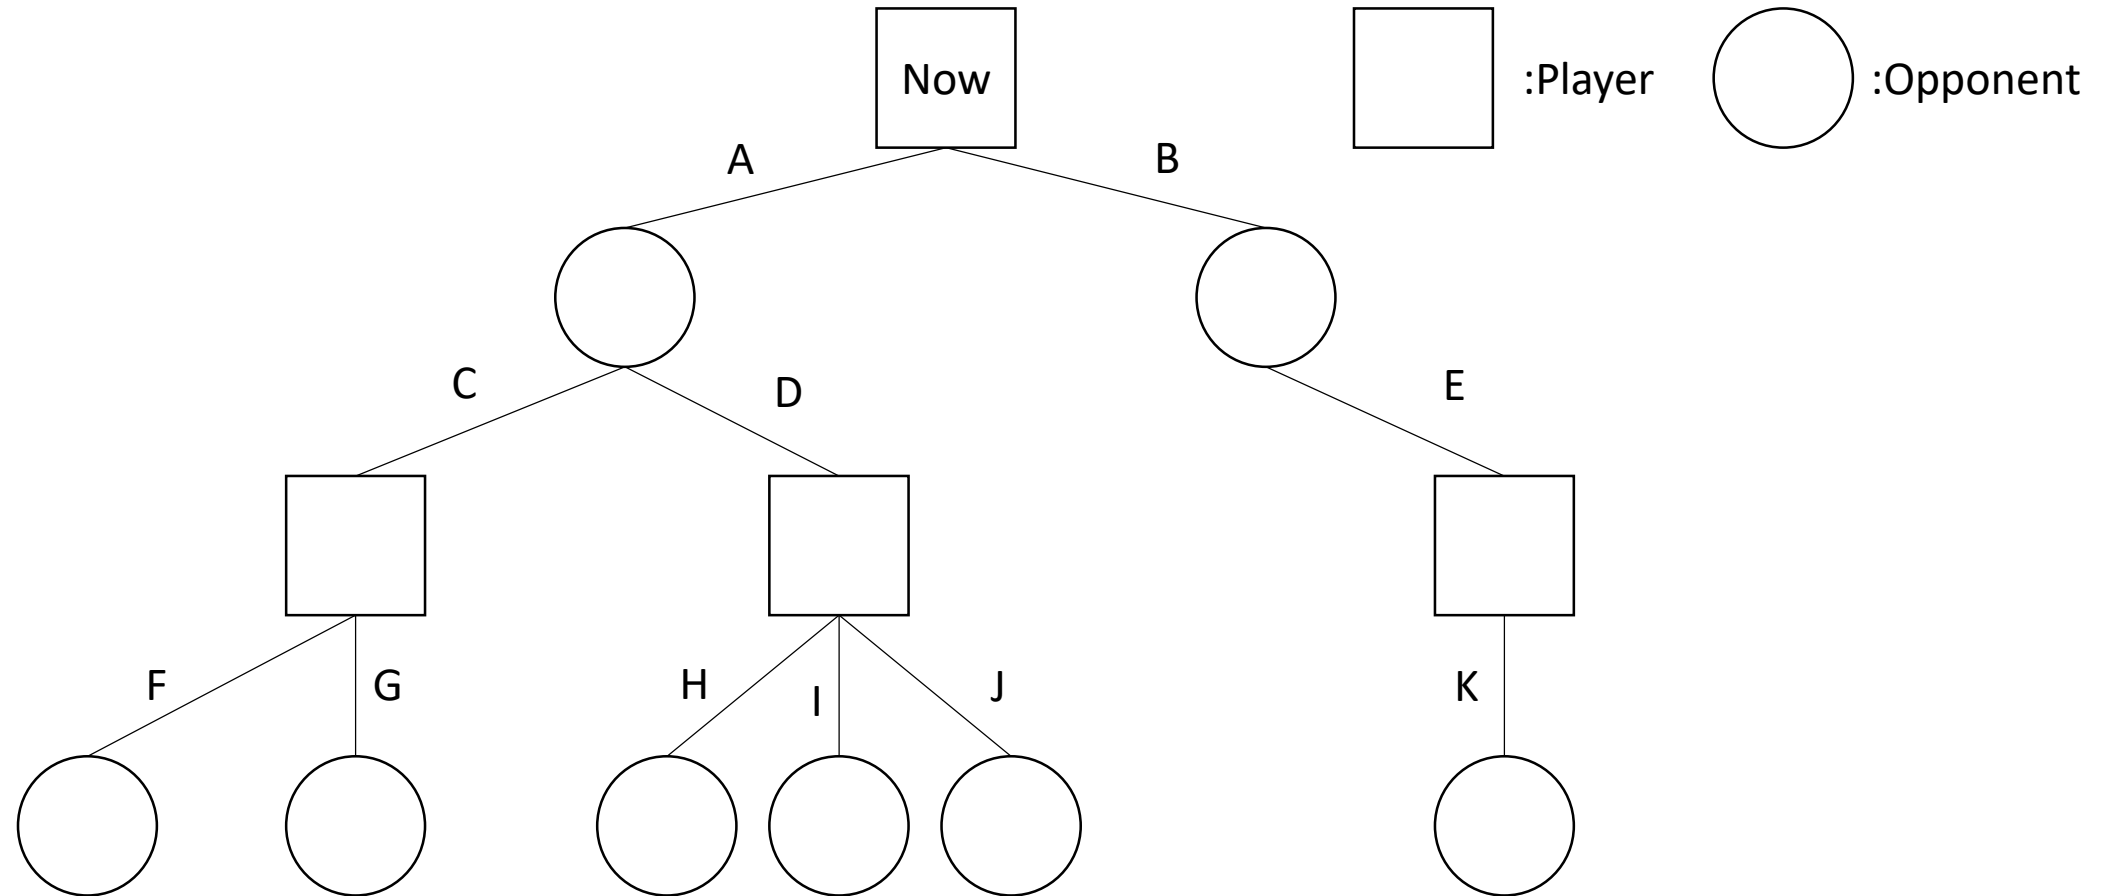
- Player tries its best to win
 - Player picks the move with the highest score
- Opponent tries its best to defeat the player
 - Opponent picks the move with the lowest “player’s value function” score
 - That is, opponent tends to **give the player the worst board**
- **The Minimax algorithm is based on this player-opponent interaction**

Minimax Pseudocode

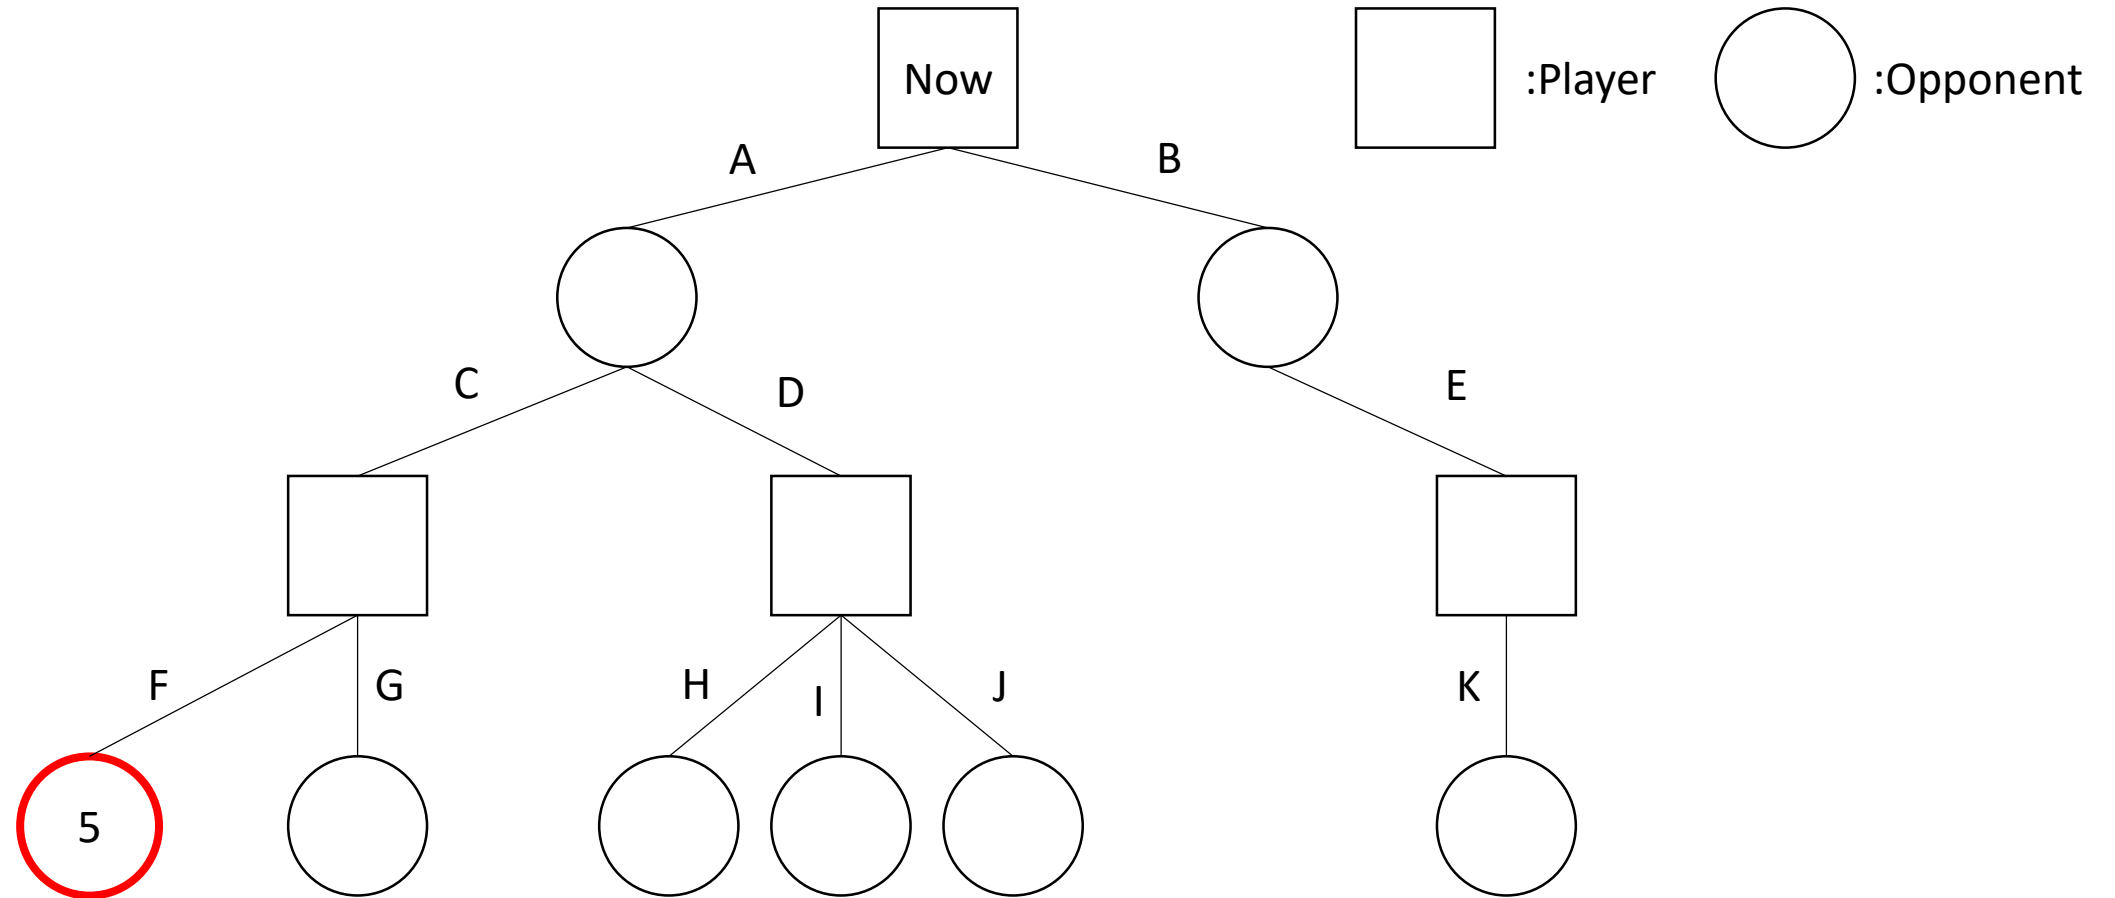
```
function minimax(node, depth, maximizingPlayer) is  
    if depth = 0 or node is a terminal node then  
        return the heuristic value of node  
    if maximizingPlayer then  
        value :=  $-\infty$   
        for each child of node do  
            value := max(value, minimax(child, depth - 1, FALSE))  
        return value  
    else (* minimizing player *)  
        value :=  $+\infty$   
        for each child of node do  
            value := min(value, minimax(child, depth - 1, TRUE))  
    return value
```

Source: <https://en.wikipedia.org/wiki/Minimax>

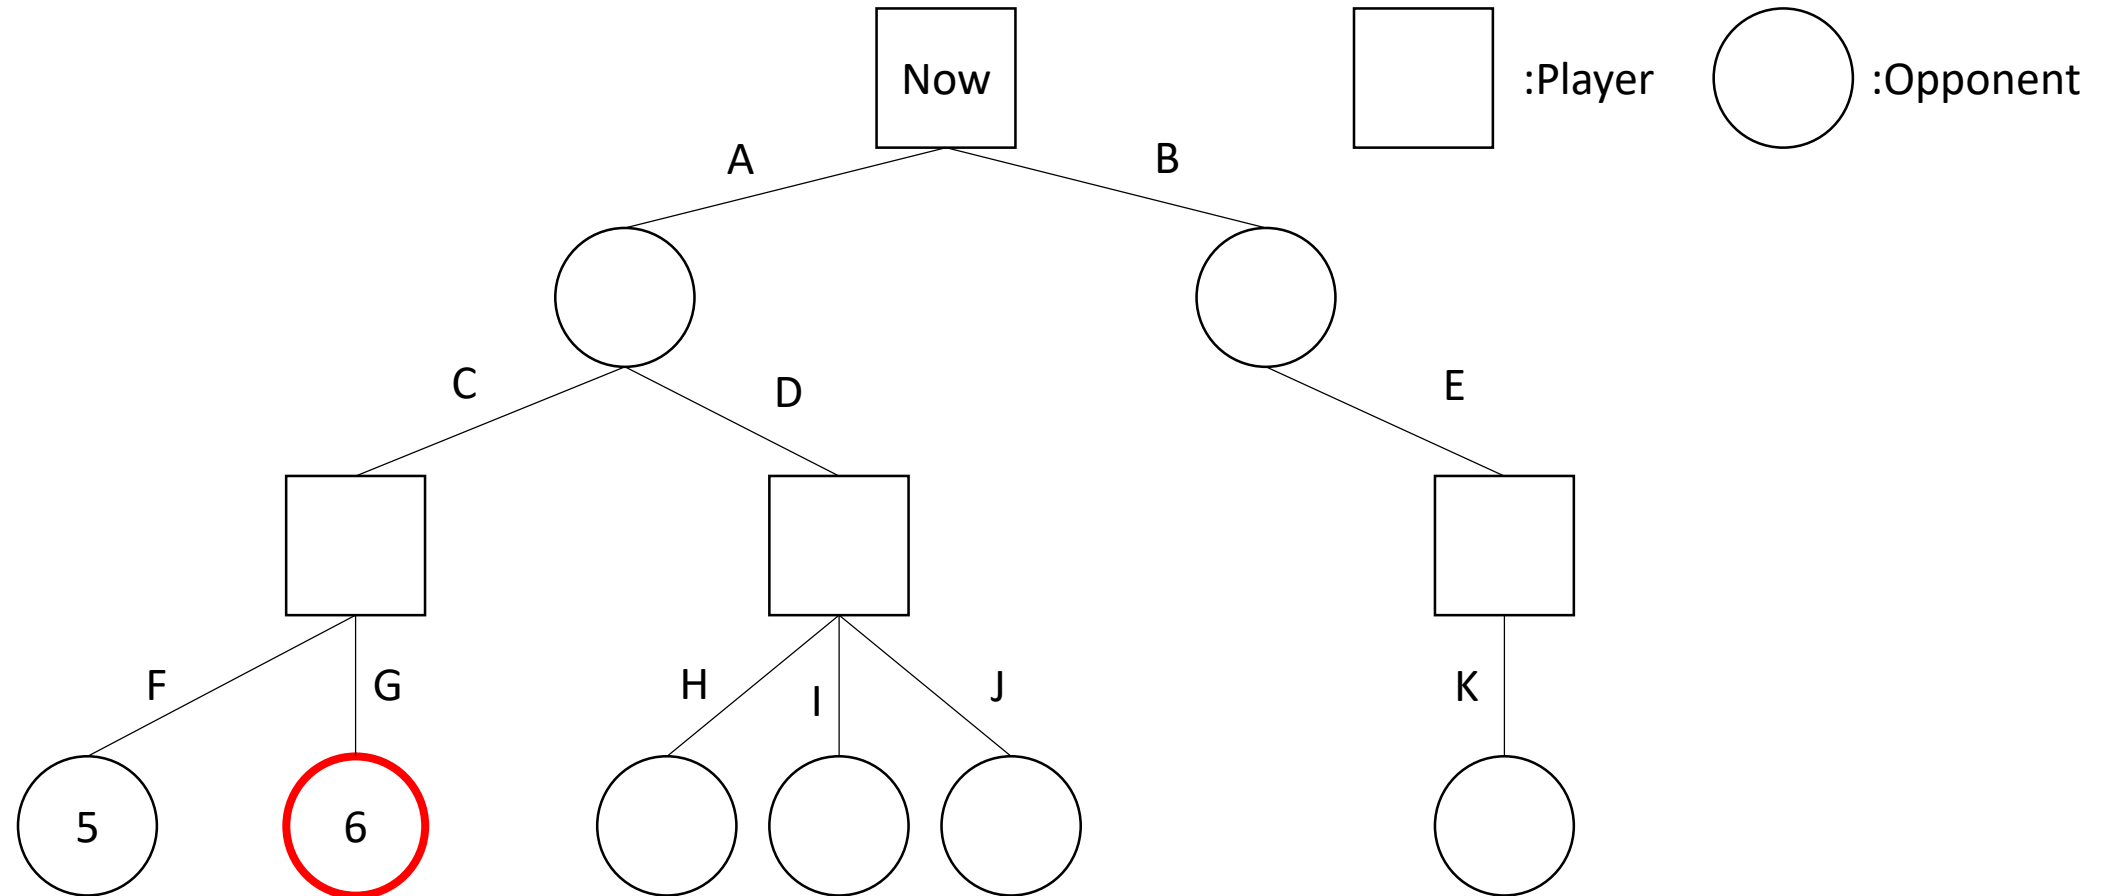
Example



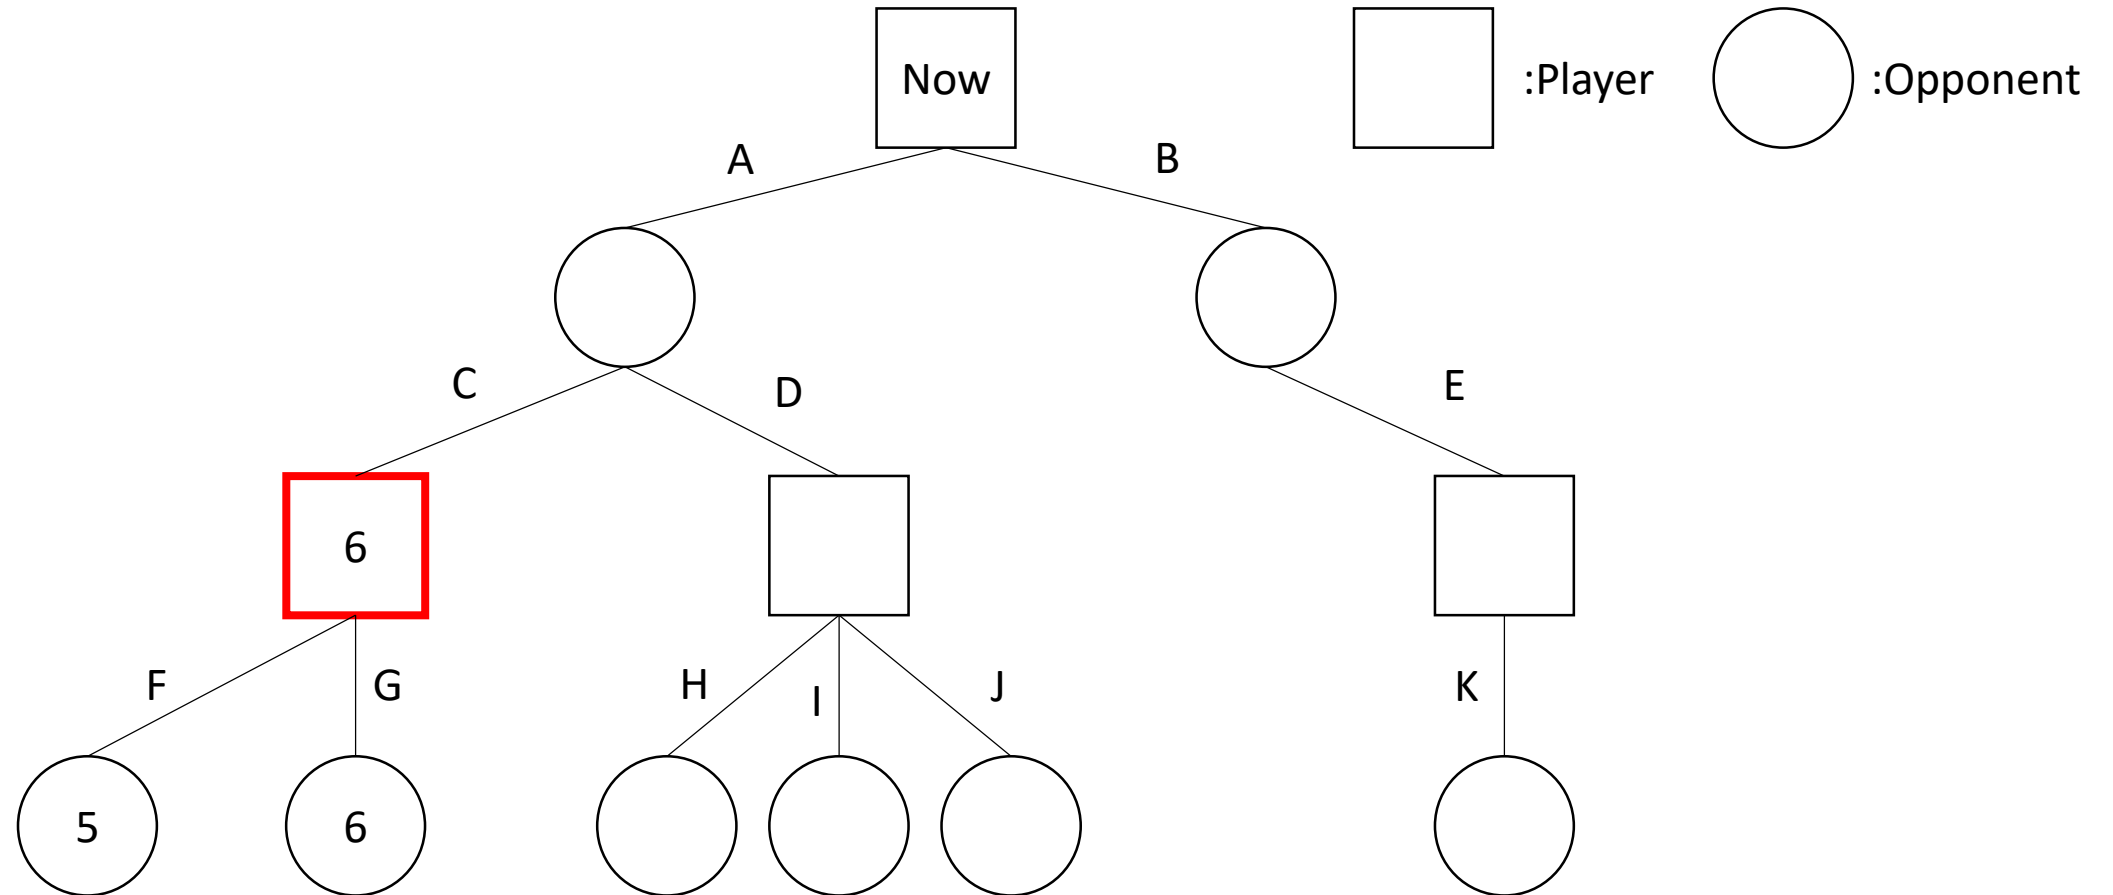
Evaluate score at leaves



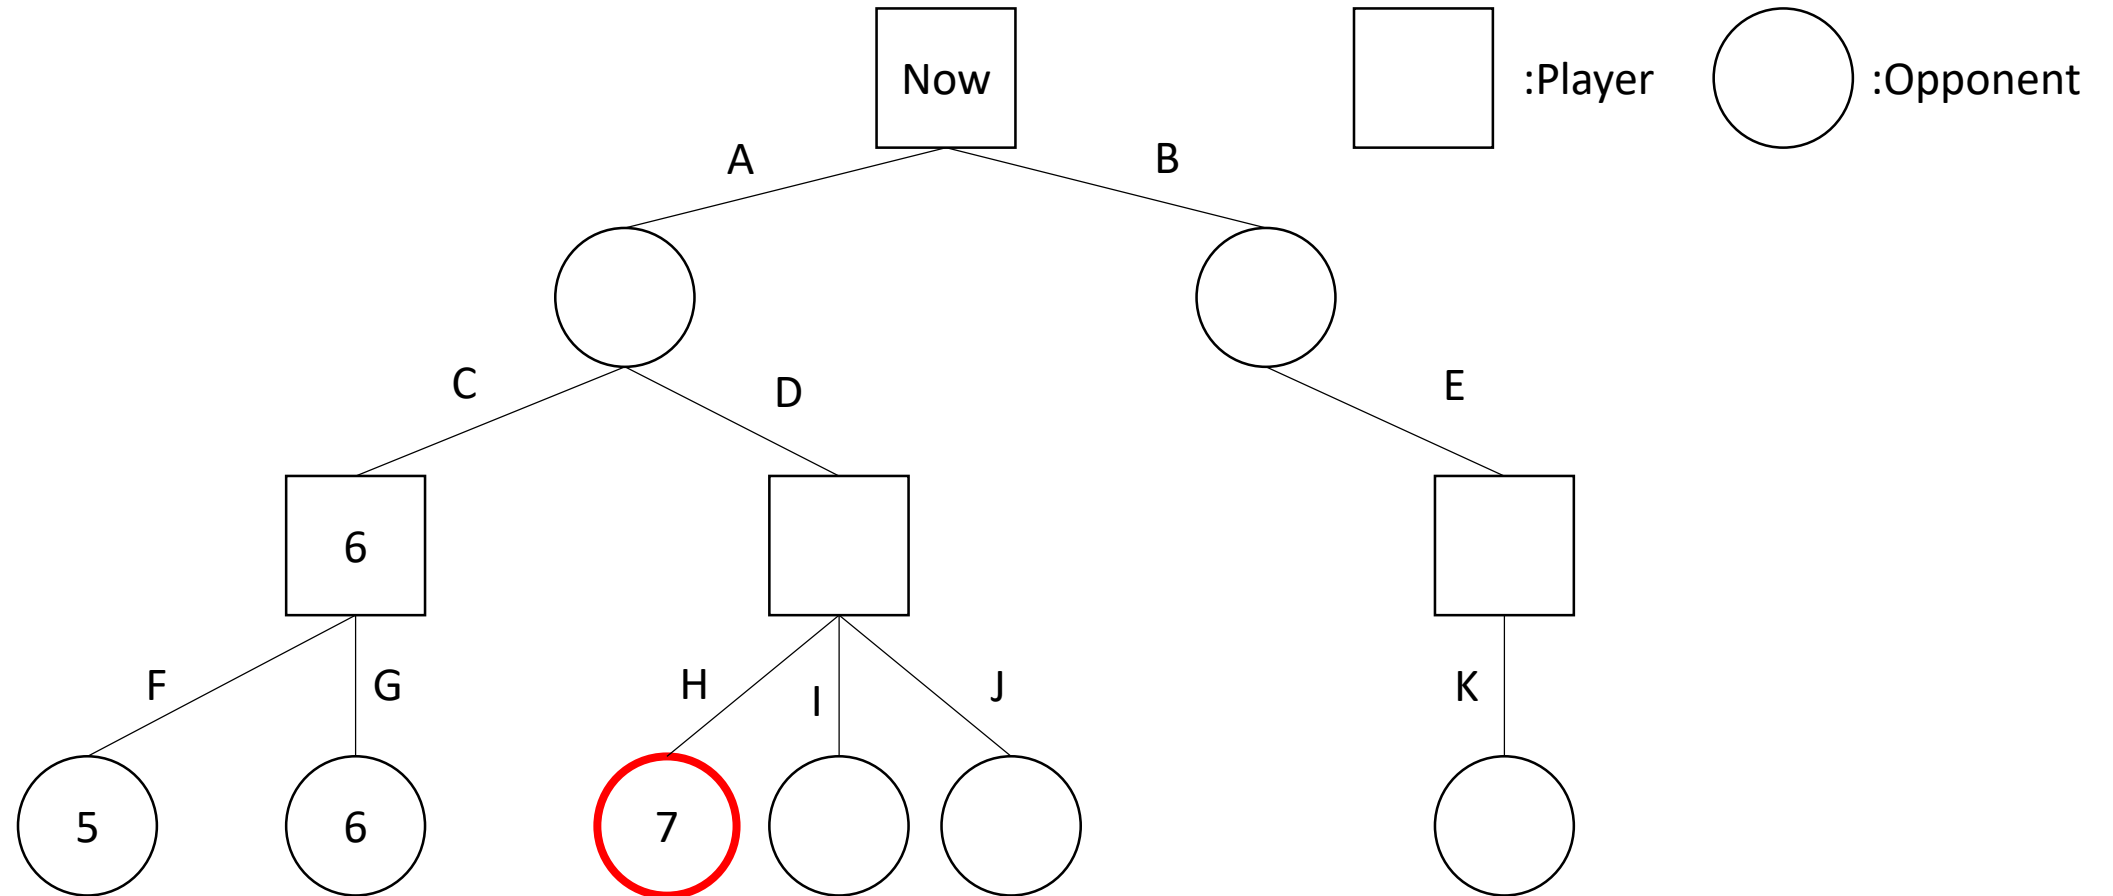
Evaluate score at leaves



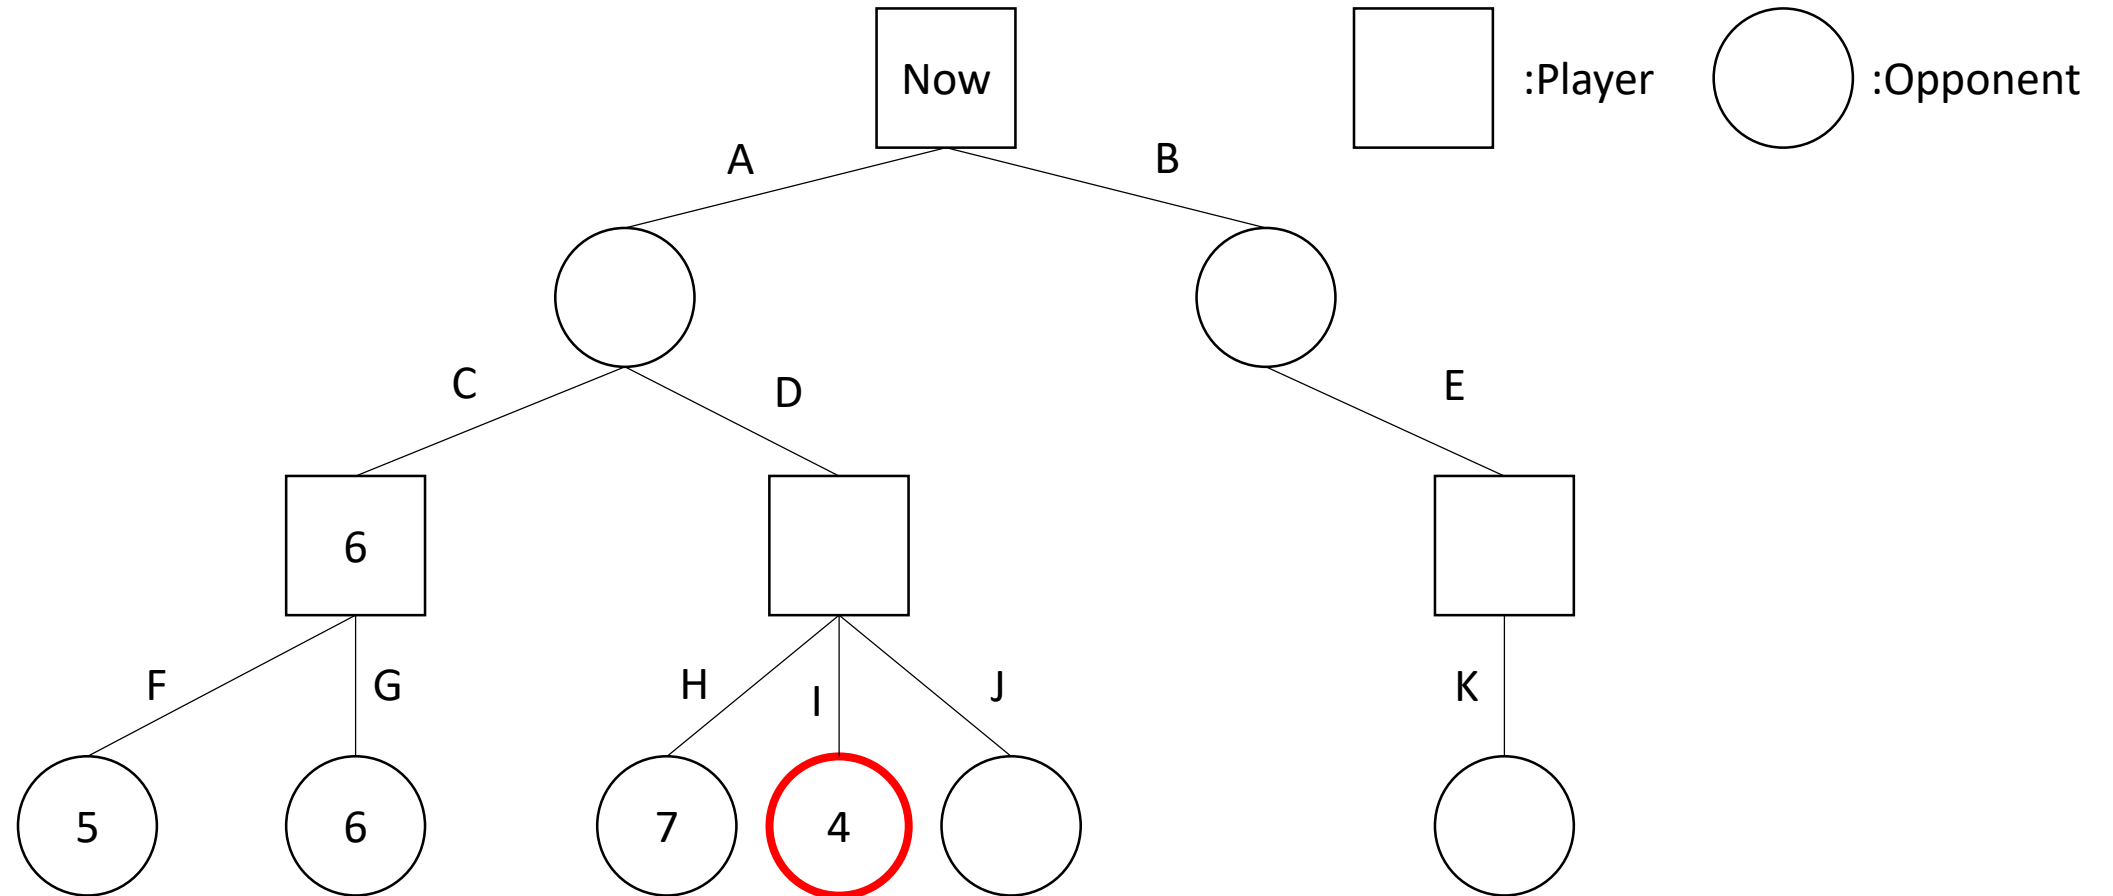
Player picks the largest score



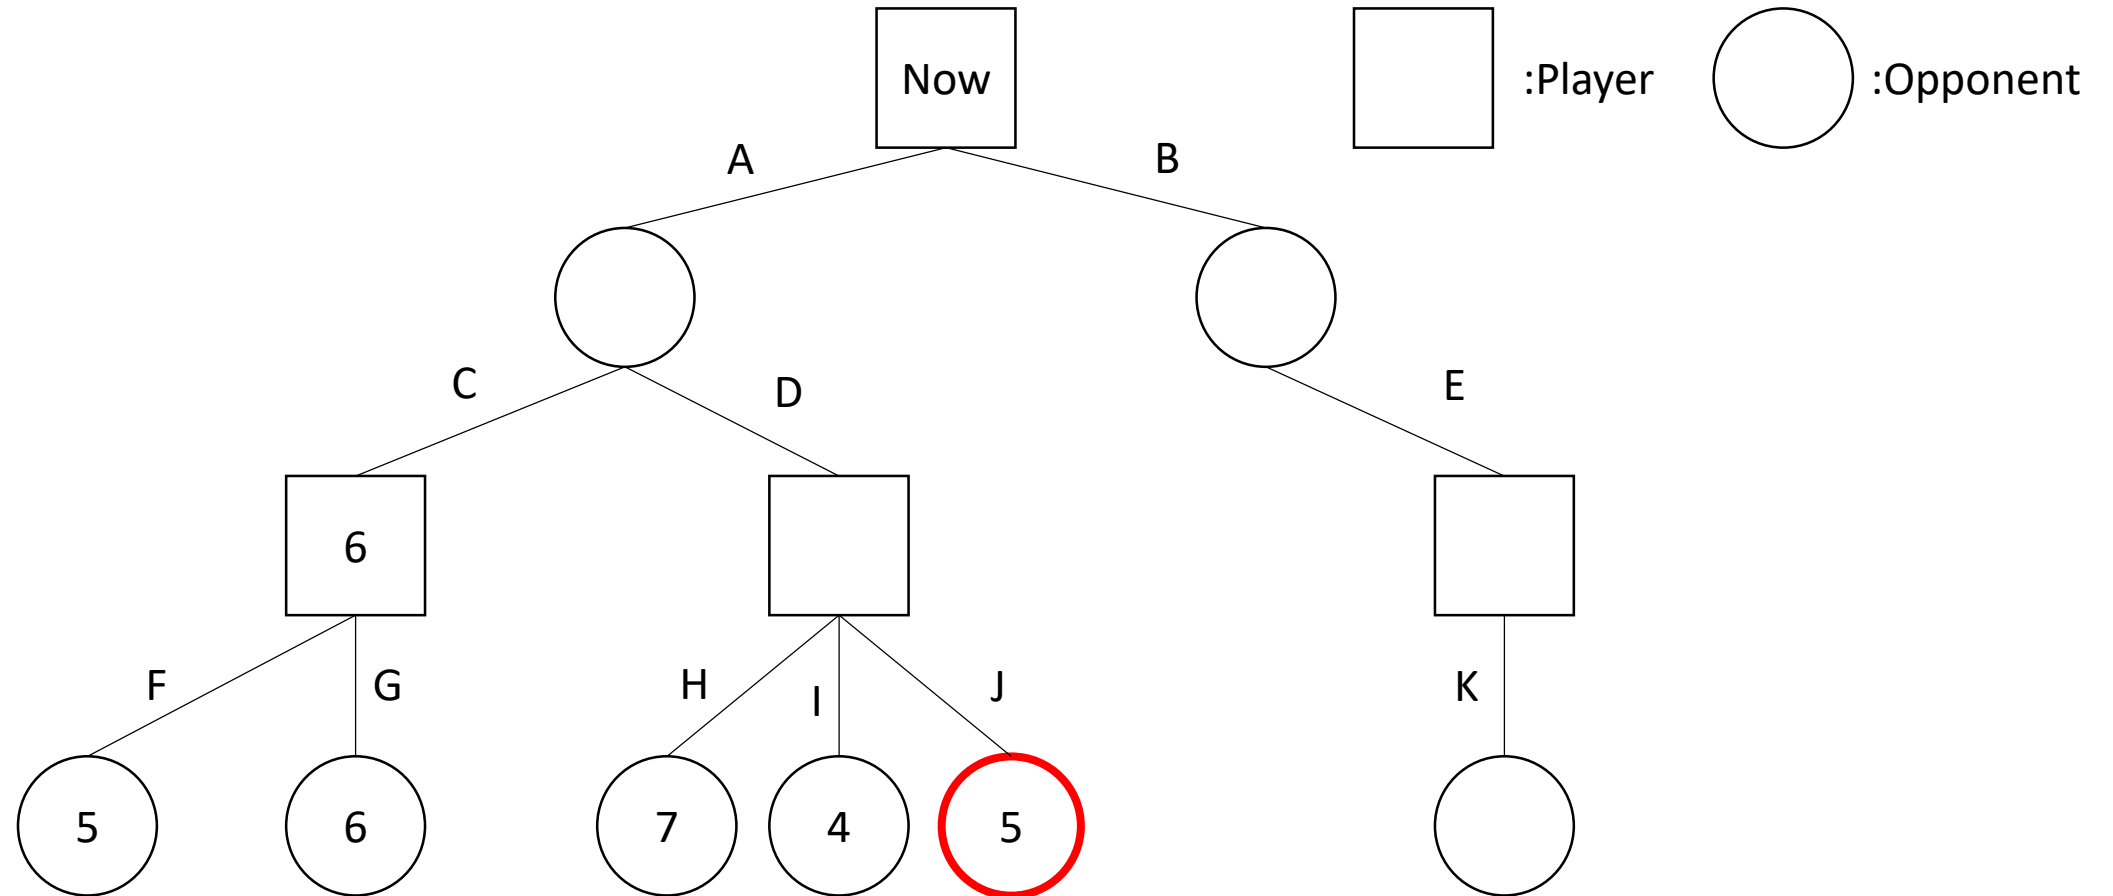
Evaluate score at leaves



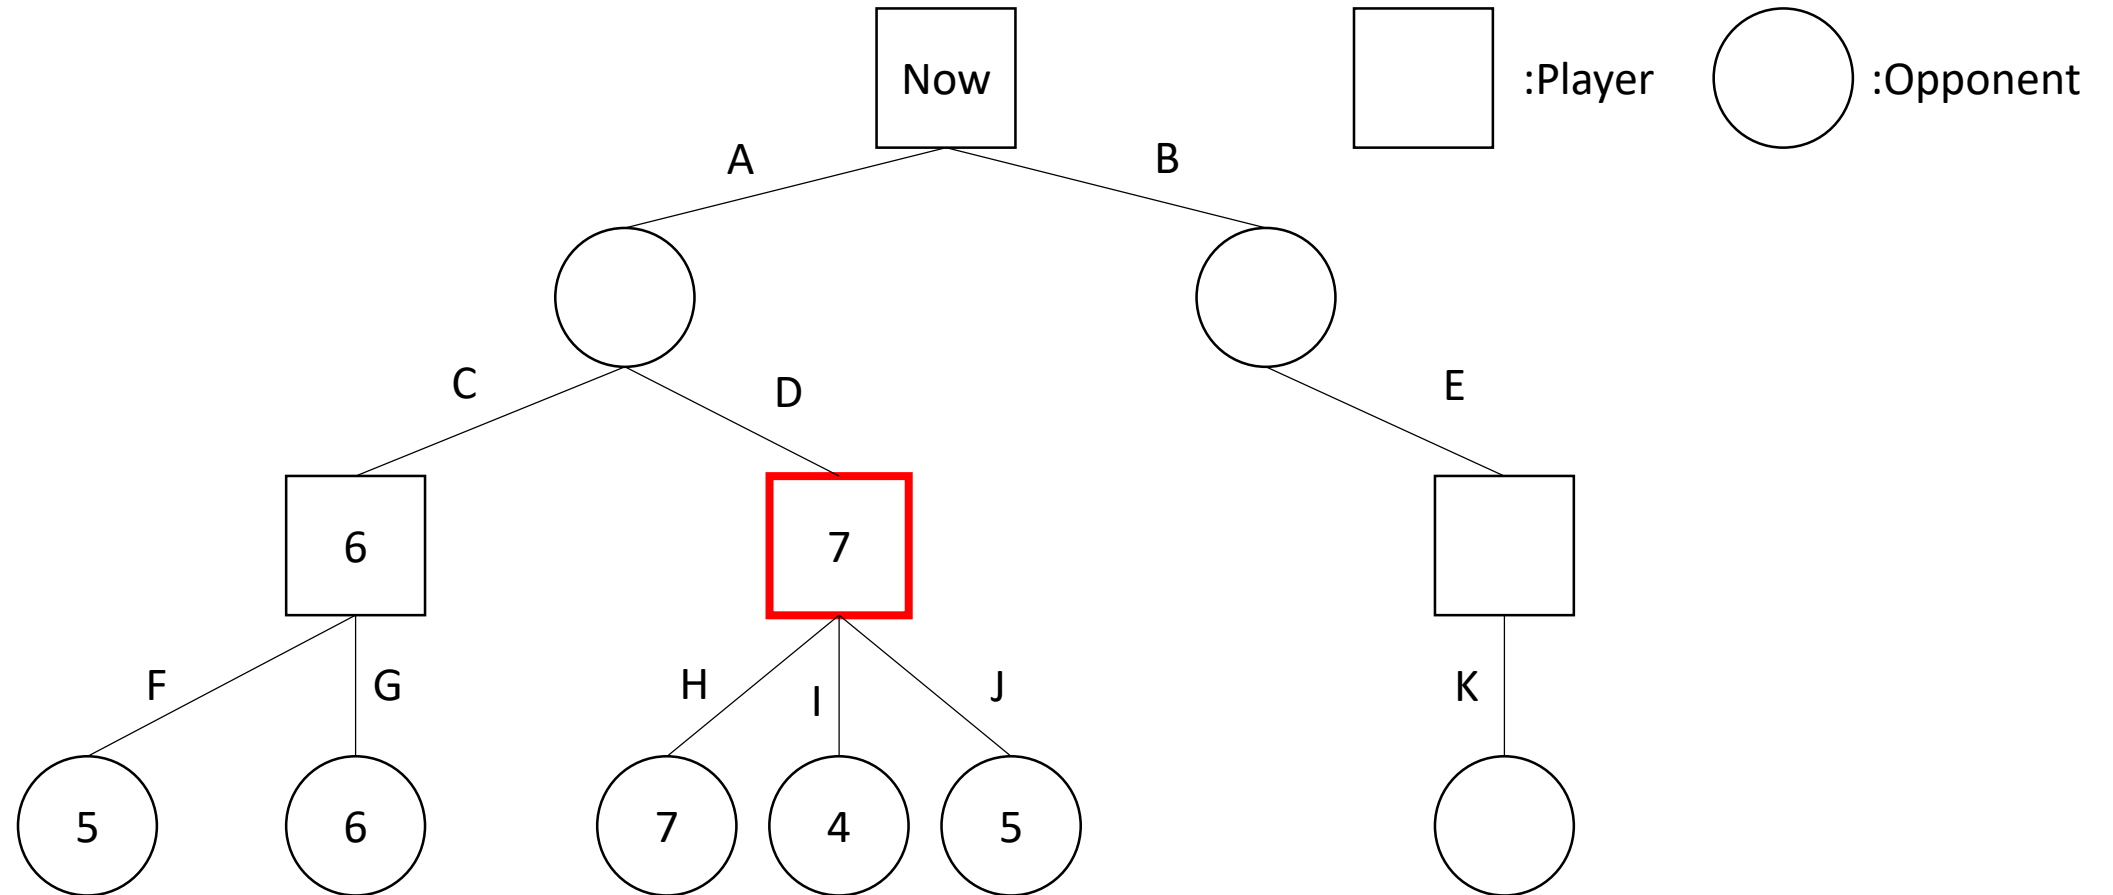
Evaluate score at leaves



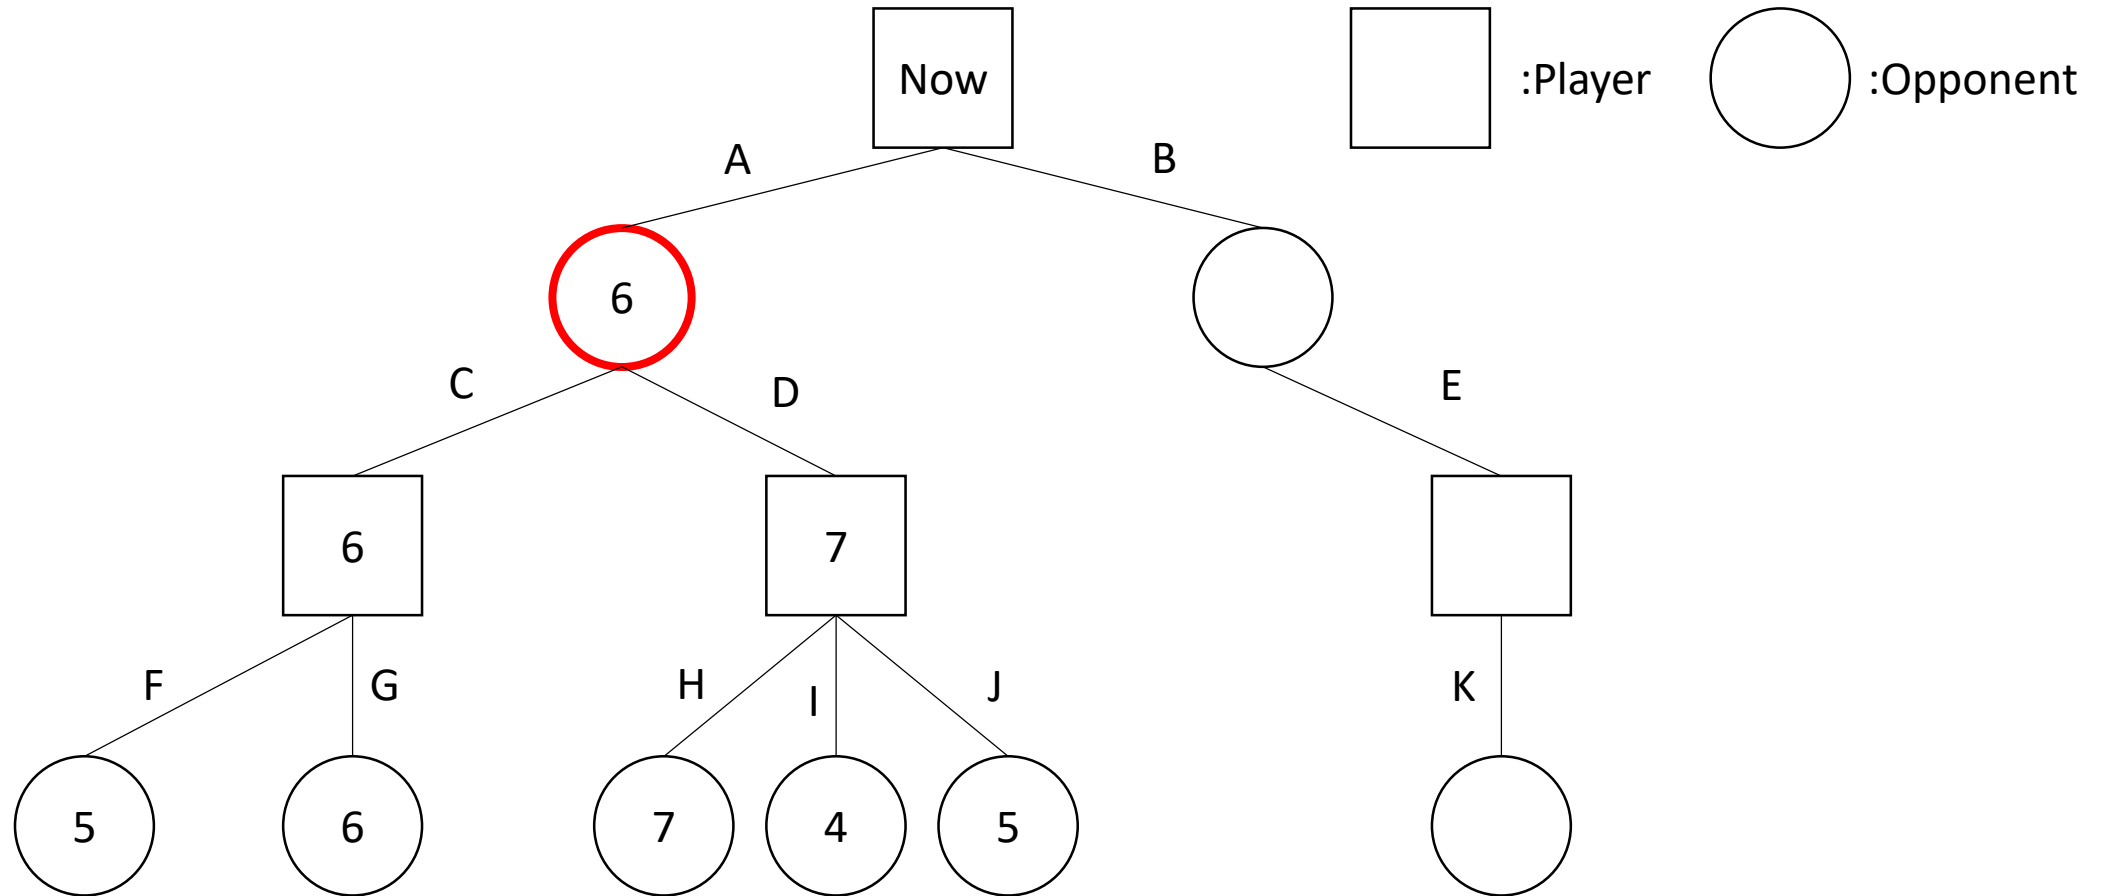
Evaluate score at leaves



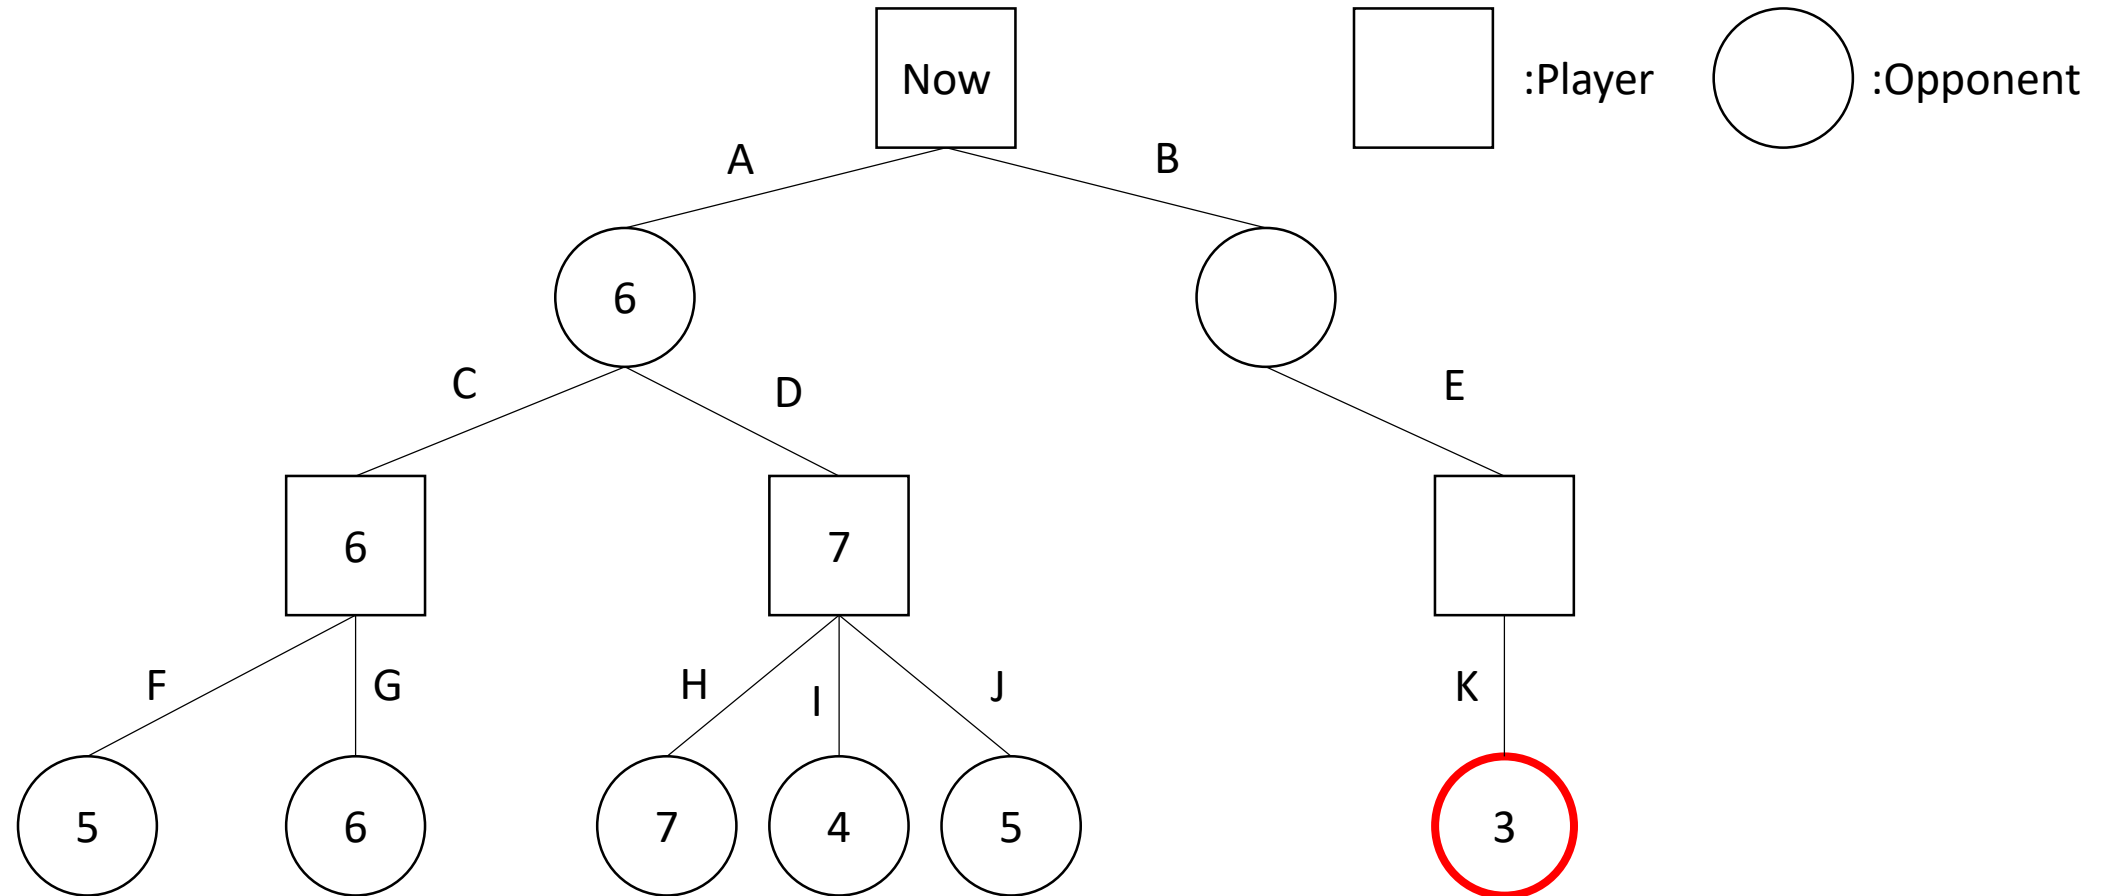
Player picks the largest score



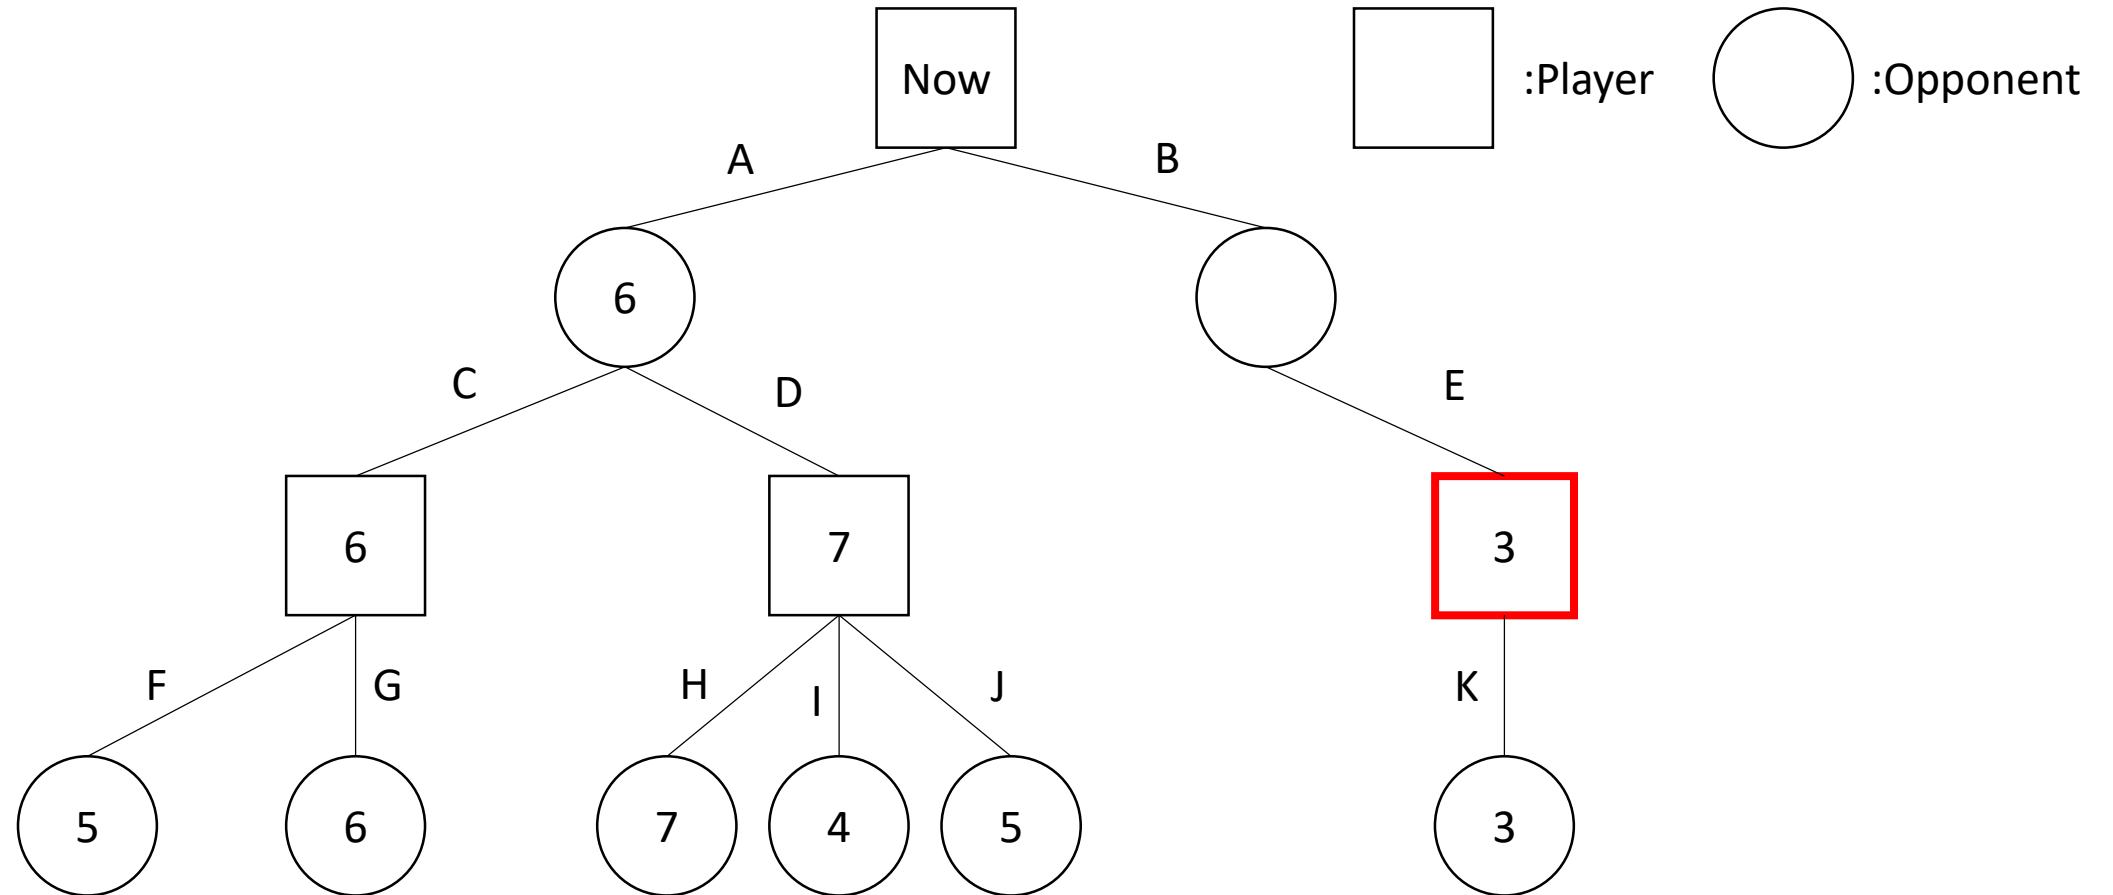
Opponent picks the smallest score



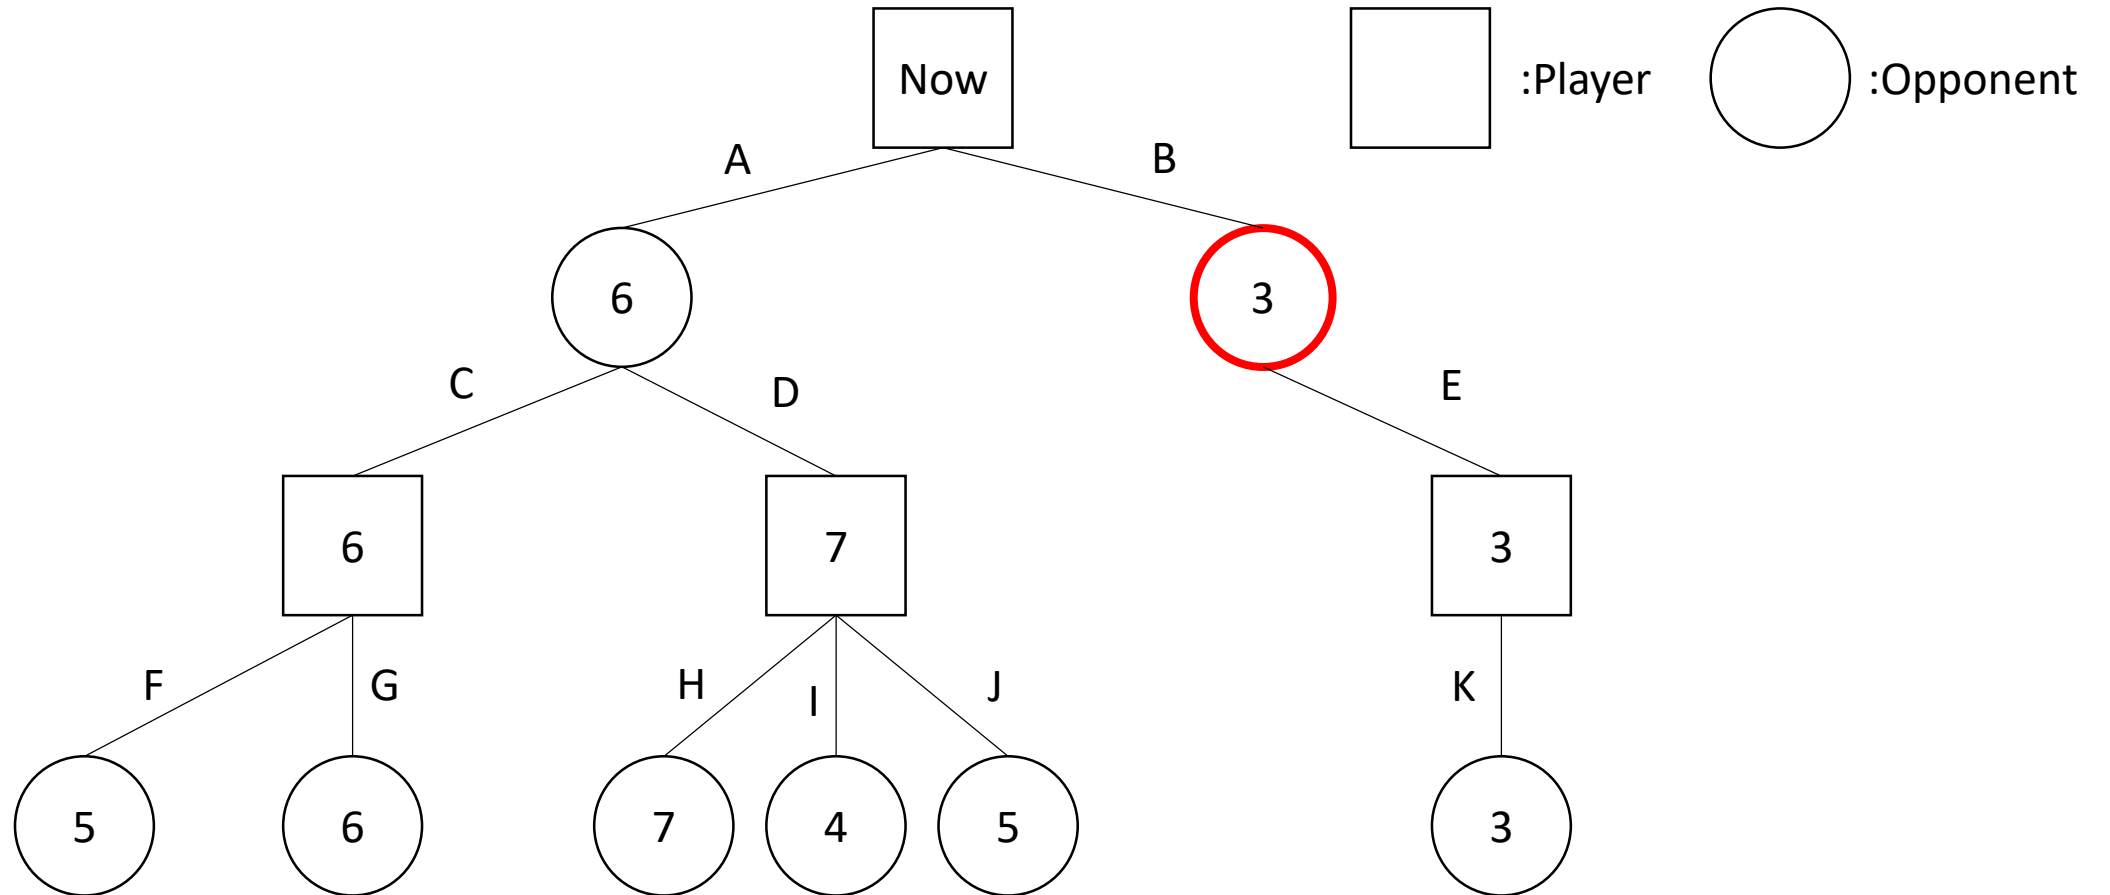
Evaluate score at leaves



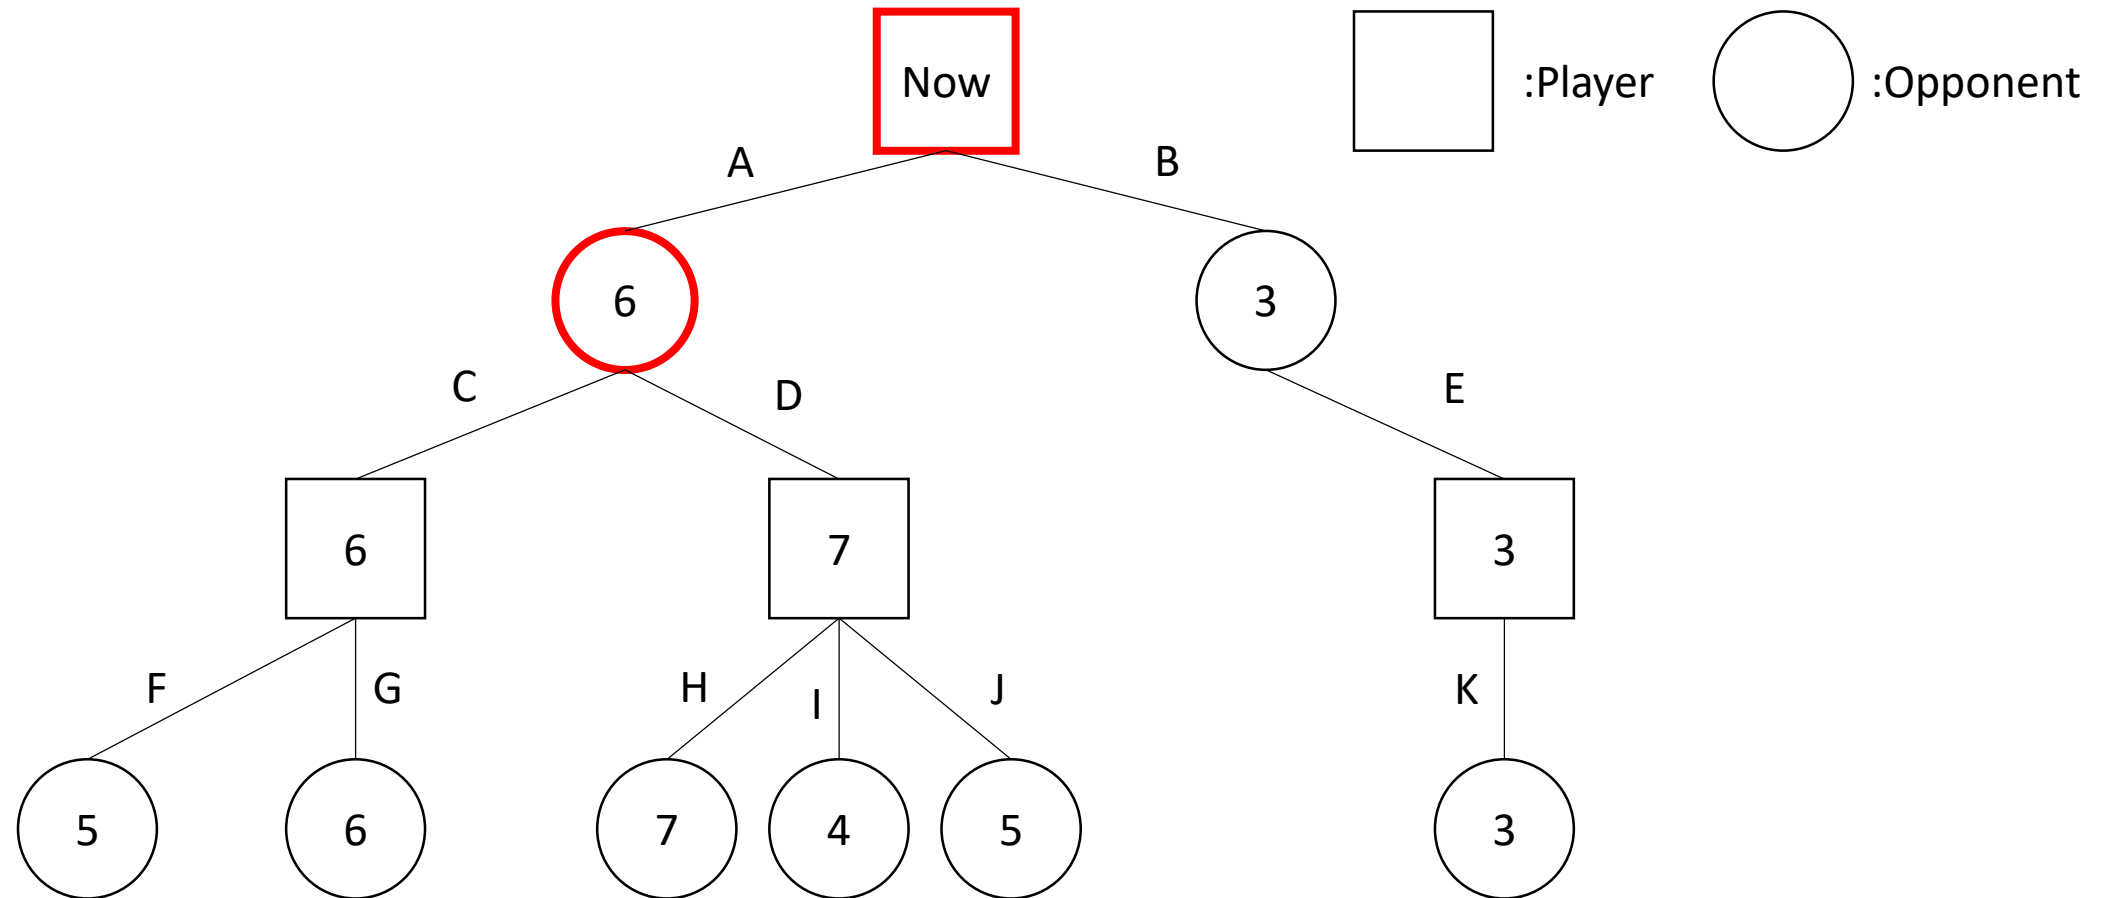
Player picks the largest score



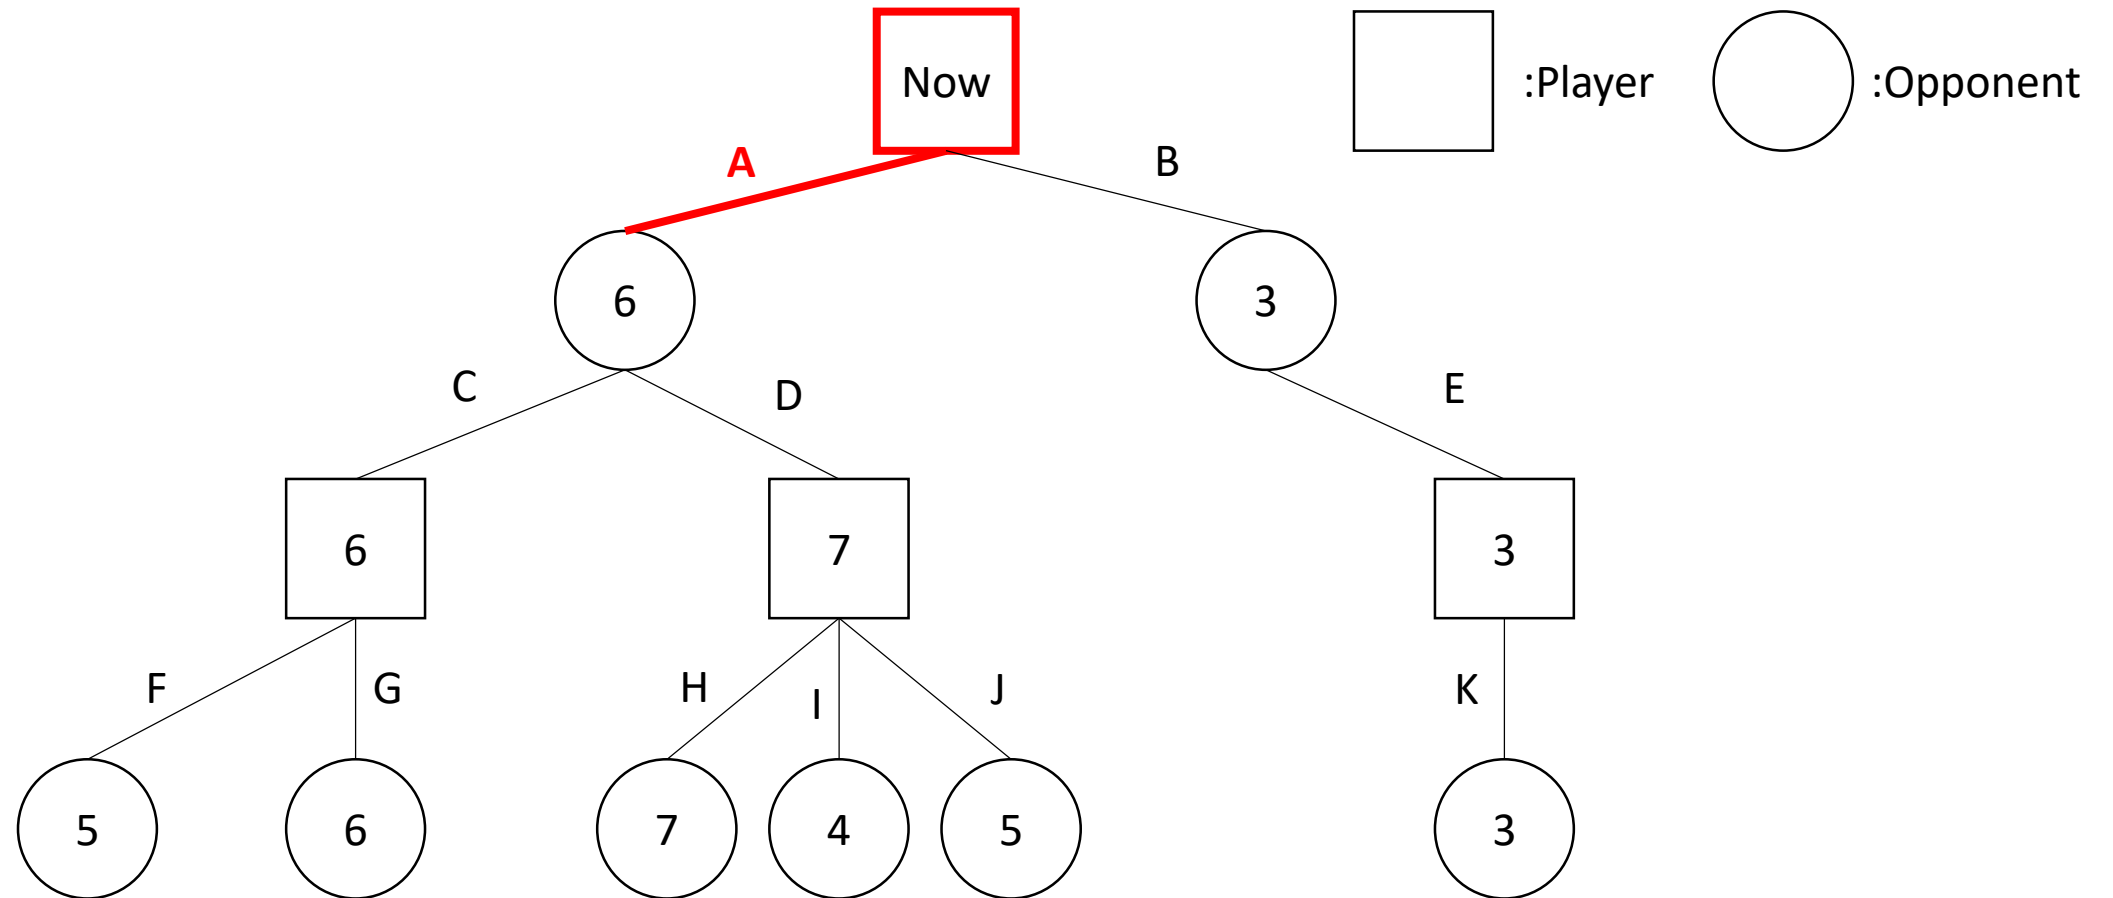
Opponent picks the smallest score



Move A has the largest score



Player picks move A to be the next move



Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. **Alpha-Beta Pruning**
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

Alpha-Beta Pruning

- By Minimax, we can simulate our opponent's moves and pick a move with minimum risk and maximum value
- Looking forward for more steps may improve the policy
- However, the size of the search tree may drastically increase with the increase of search depth


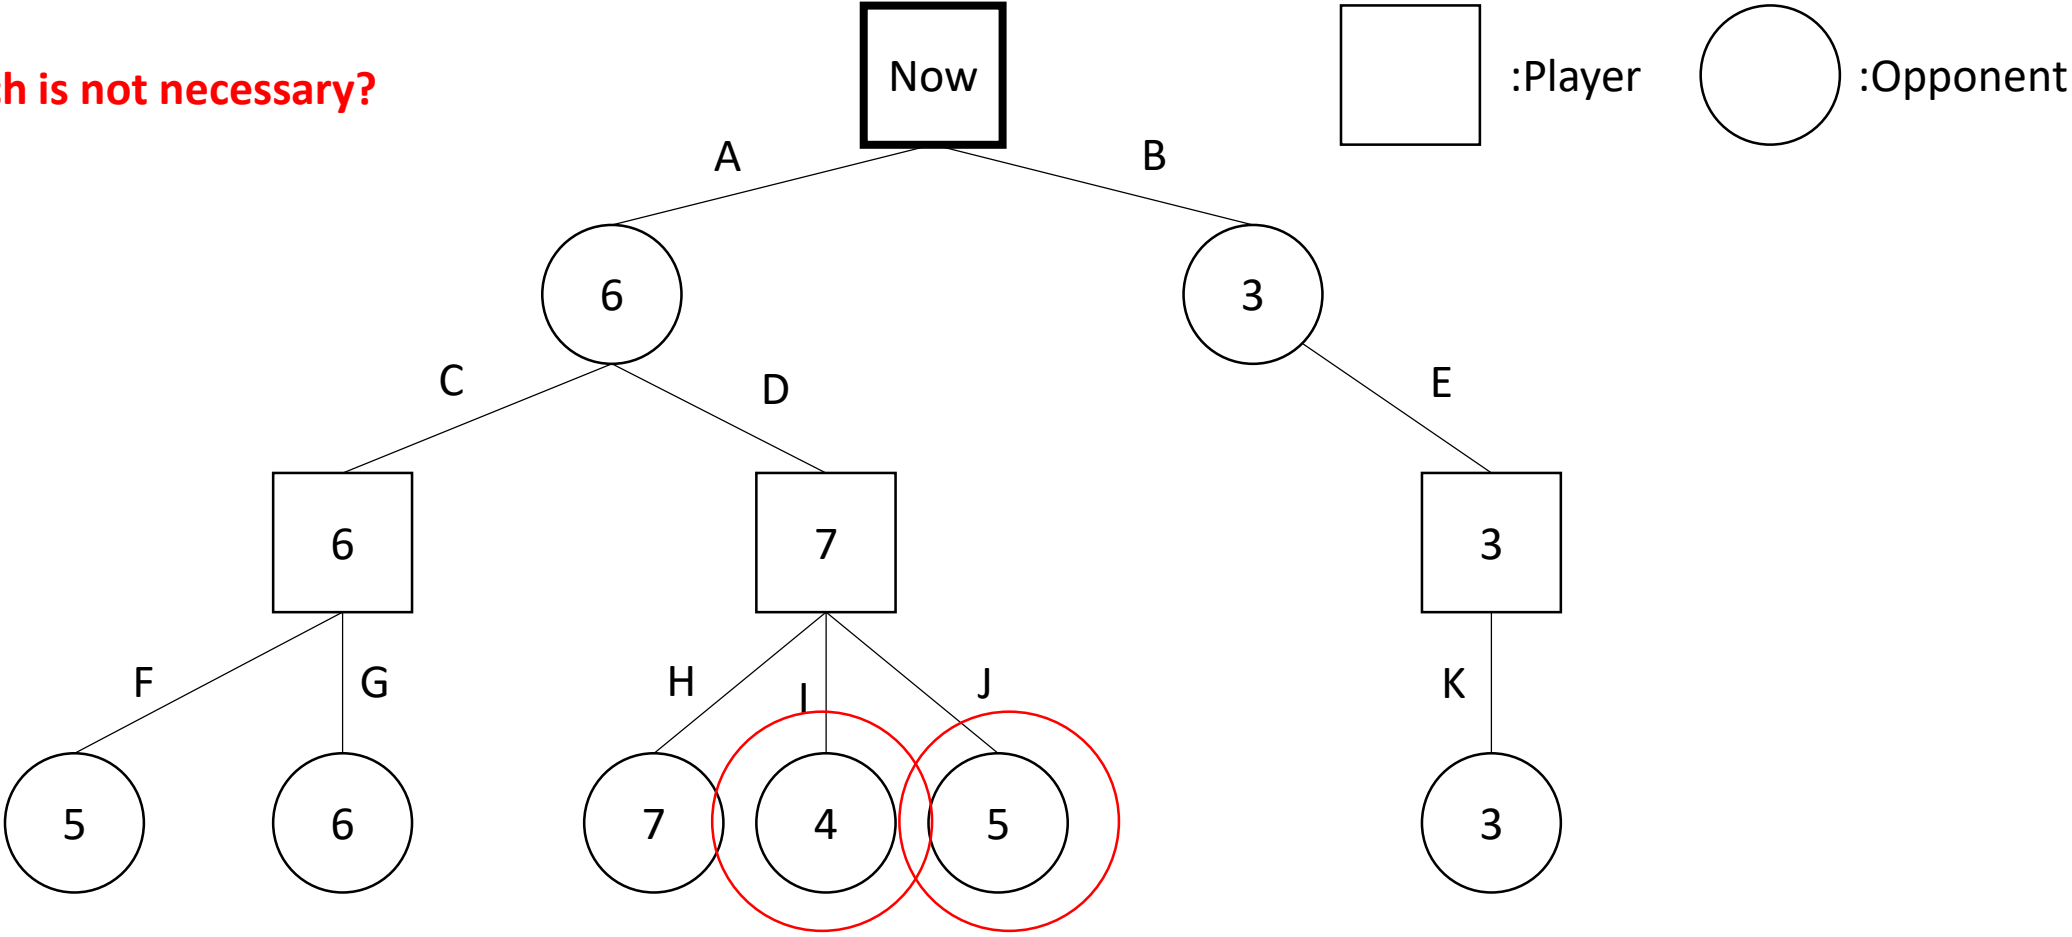


Figure 4.26 A randomly generated binary search tree

Alpha-Beta Pruning

- Since we only have limited time, if we hope to increase search depth, we must optimize the search process
- There are many branches in the minimax process which is not related to the result
- We can try to “prune” these branches to improve efficiency
- The Alpha-Beta Pruning is the improved version of Minimax method which eliminates some unnecessary branches

Which branch is not necessary?



Alpha-Beta Pruning Pseudocode

```
function alphabeta(node, depth,  $\alpha$ ,  $\beta$ , maximizingPlayer) is
    if depth = 0 or node is a terminal node then
        return the heuristic value of node
    if maximizingPlayer then
        value :=  $-\infty$ 
        for each child of node do
            value := max(value, alphabeta(child, depth - 1,  $\alpha$ ,  $\beta$ , FALSE))
             $\alpha$  := max( $\alpha$ , value)
            if  $\alpha \geq \beta$  then
                break (*  $\beta$  cutoff *)
        return value
    else
        value :=  $+\infty$ 
        for each child of node do
            value := min(value, alphabeta(child, depth - 1,  $\alpha$ ,  $\beta$ , TRUE))
             $\beta$  := min( $\beta$ , value)
            if  $\beta \leq \alpha$  then
                break (*  $\alpha$  cutoff *)
        return value
```

Source: https://en.wikipedia.org/wiki/Alpha%E2%80%93beta_pruning

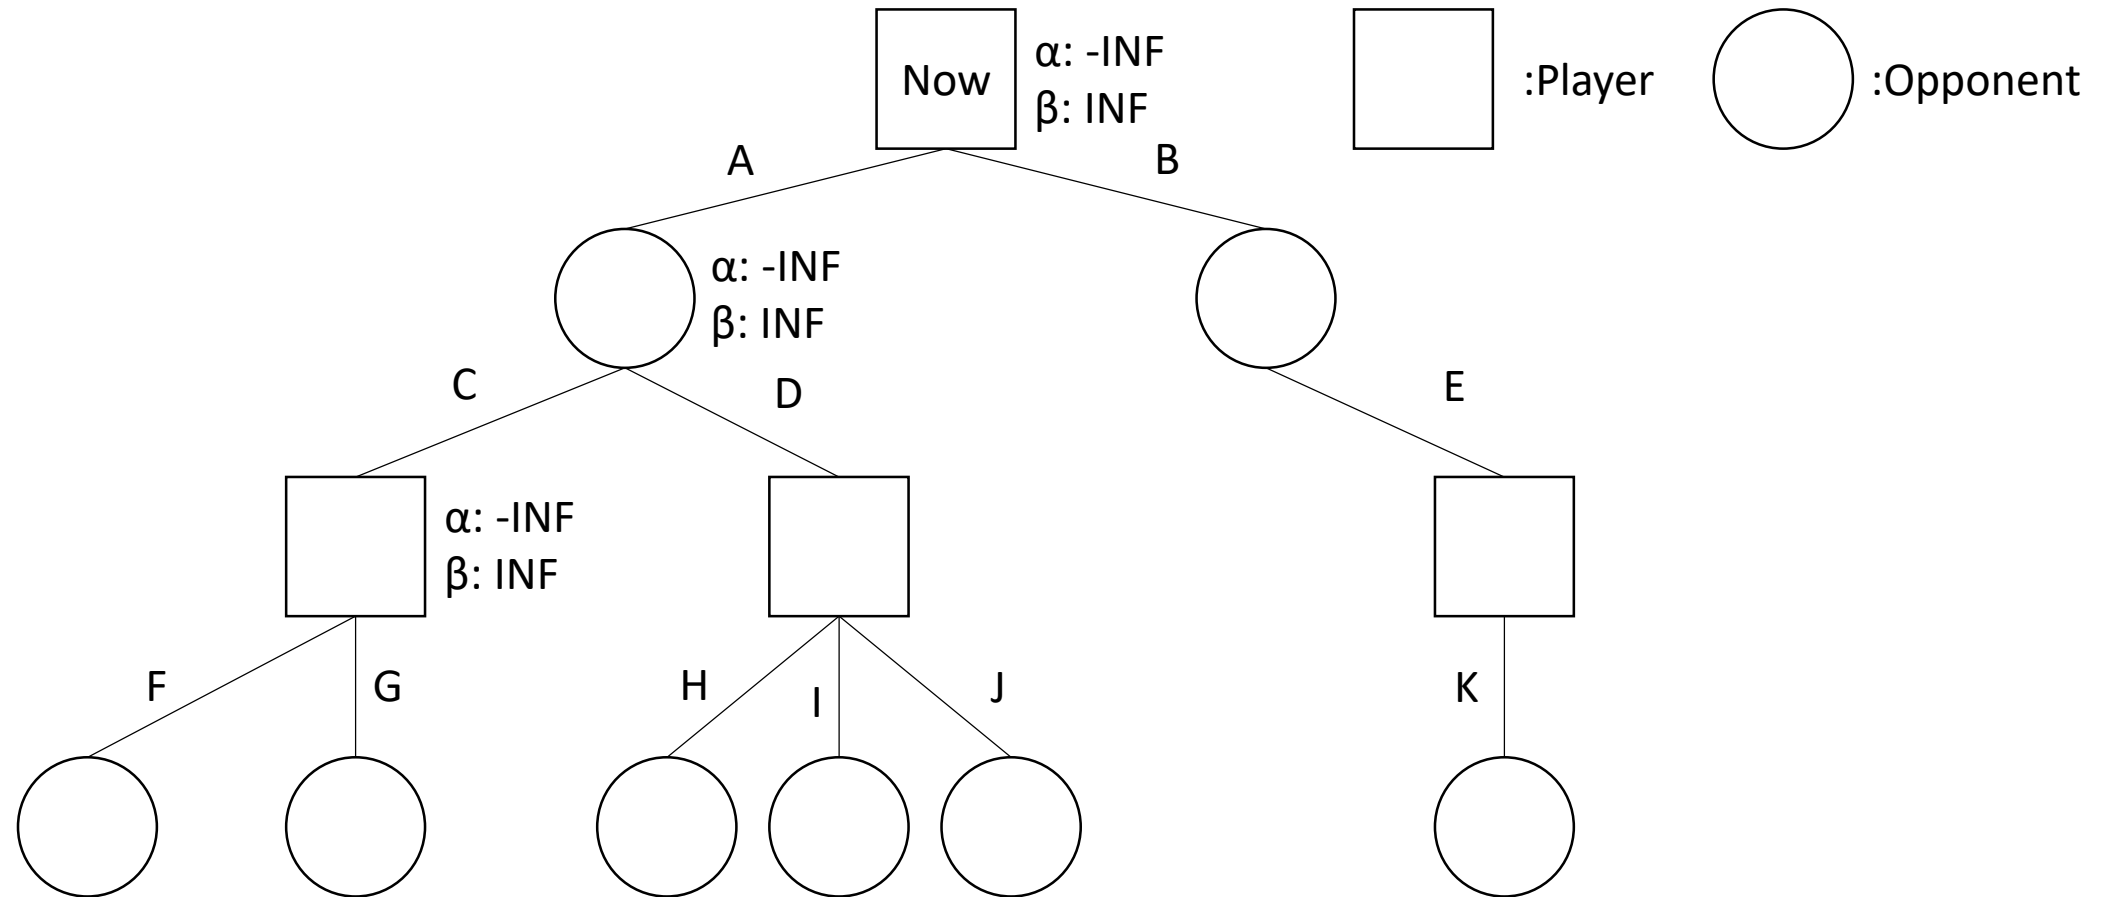
Alpha-Beta Pruning

- Alpha: the maximum score that the player is assured of in the current search process
- Beta: the minimum score that the opponent is assured of in the current search process

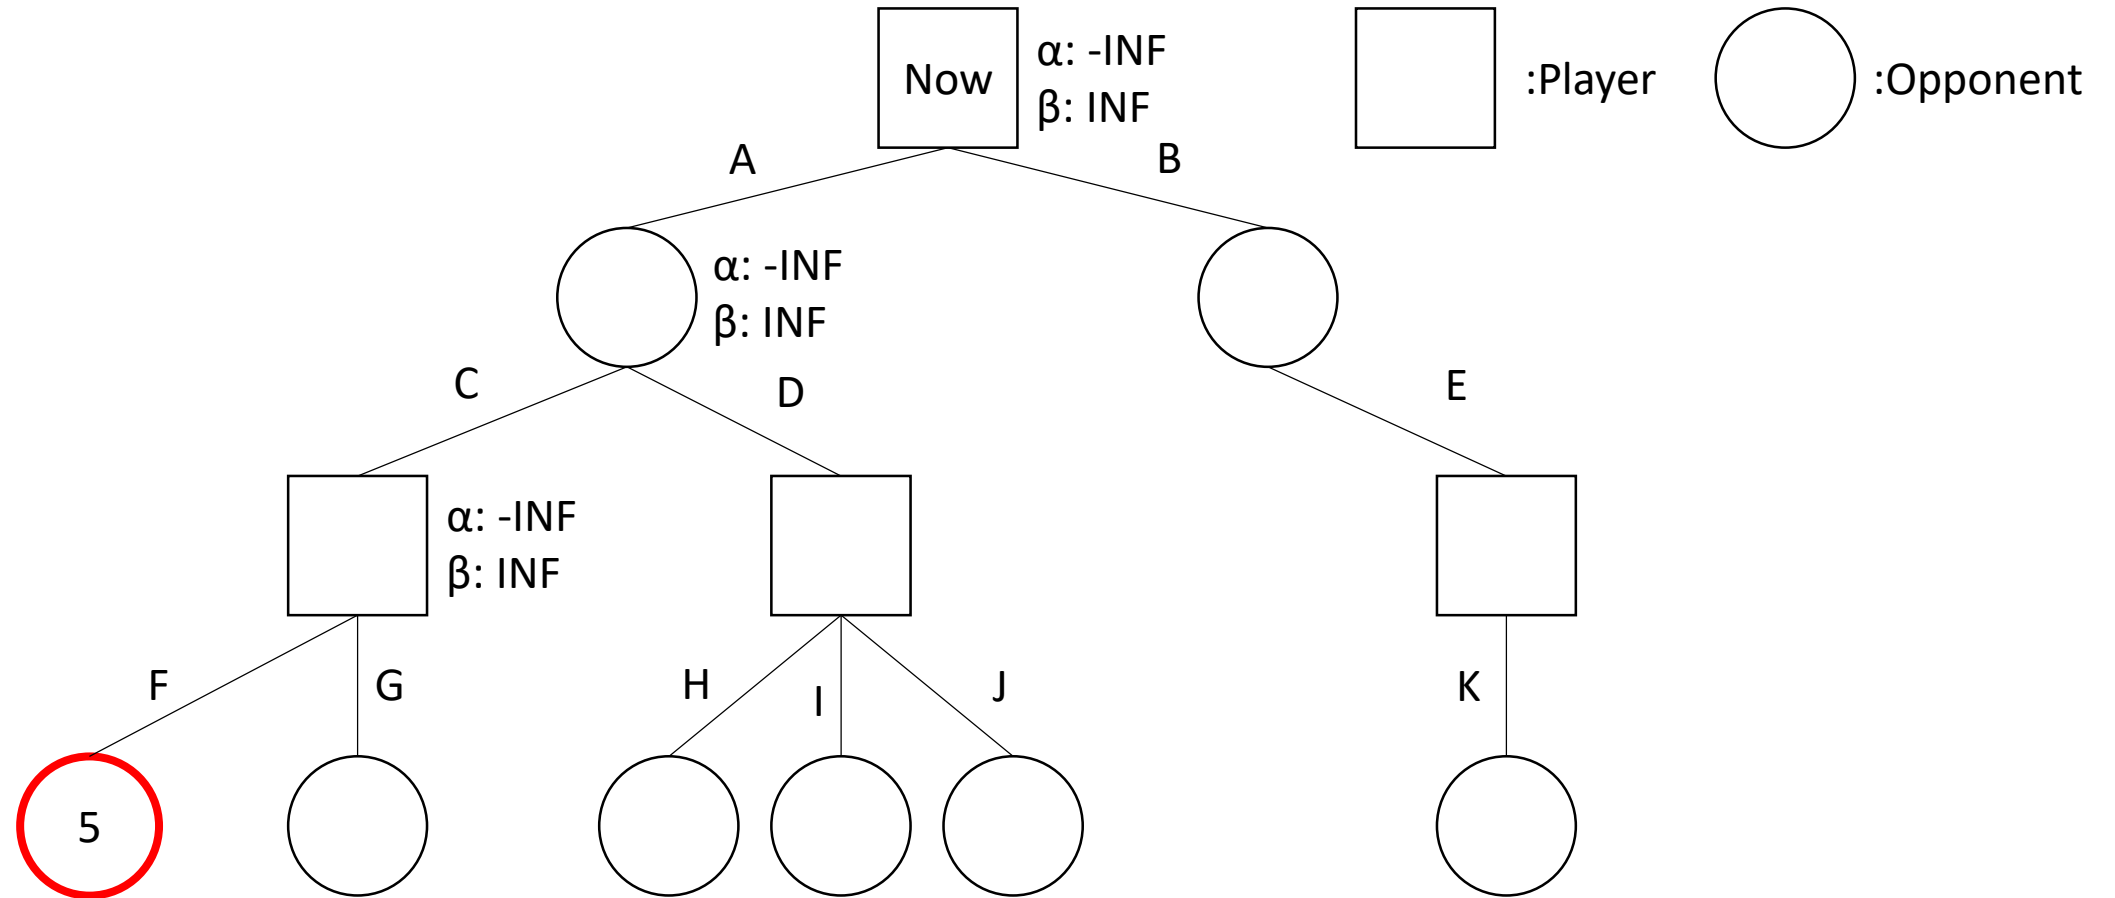
Alpha-Beta Pruning

- If $\alpha \geq \beta$ on a player node, we can stop to search on this branch
- In this situation, the player will return a value $\geq \beta$ on this branch
- However, the opponent already has a better choice (β)
- Thus, no matter the later discovered value on this branch, the opponent will not pick this branch
- We can “prune” this branch since it will not affect the result
- We can also stop to search if $\beta \leq \alpha$ on an opponent node

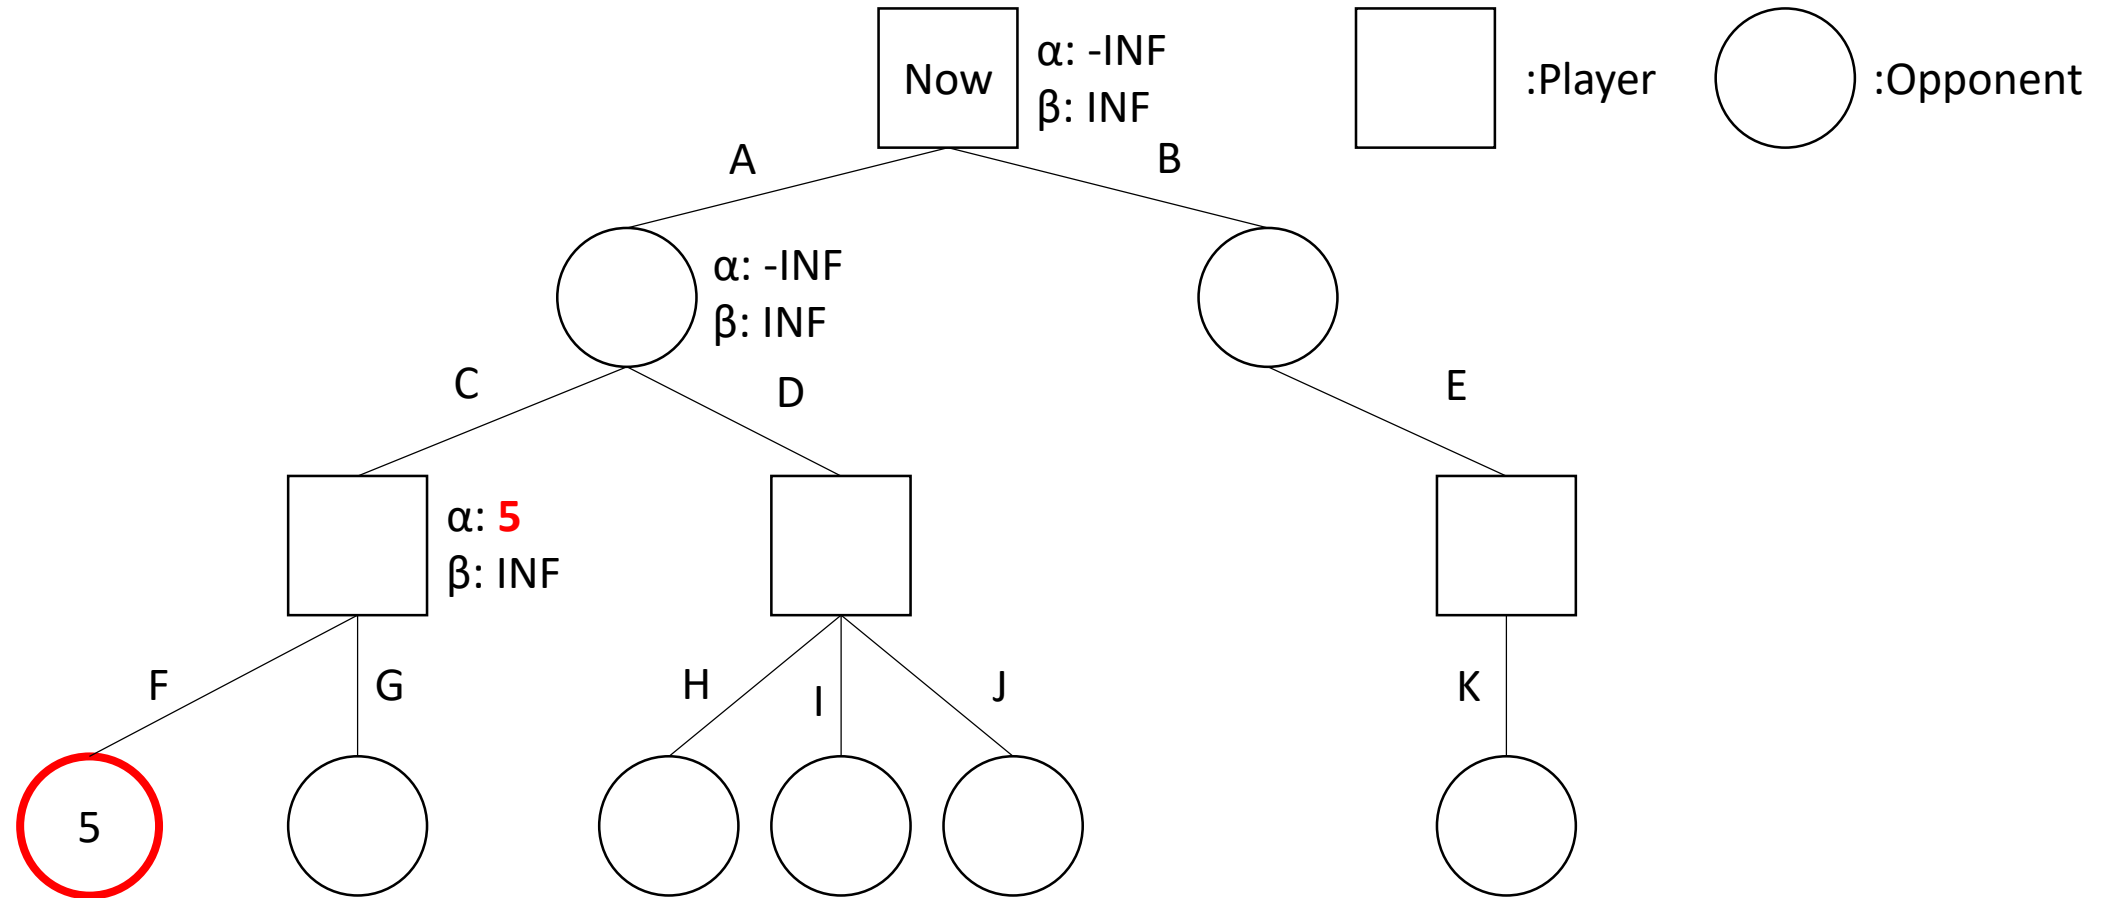
Example



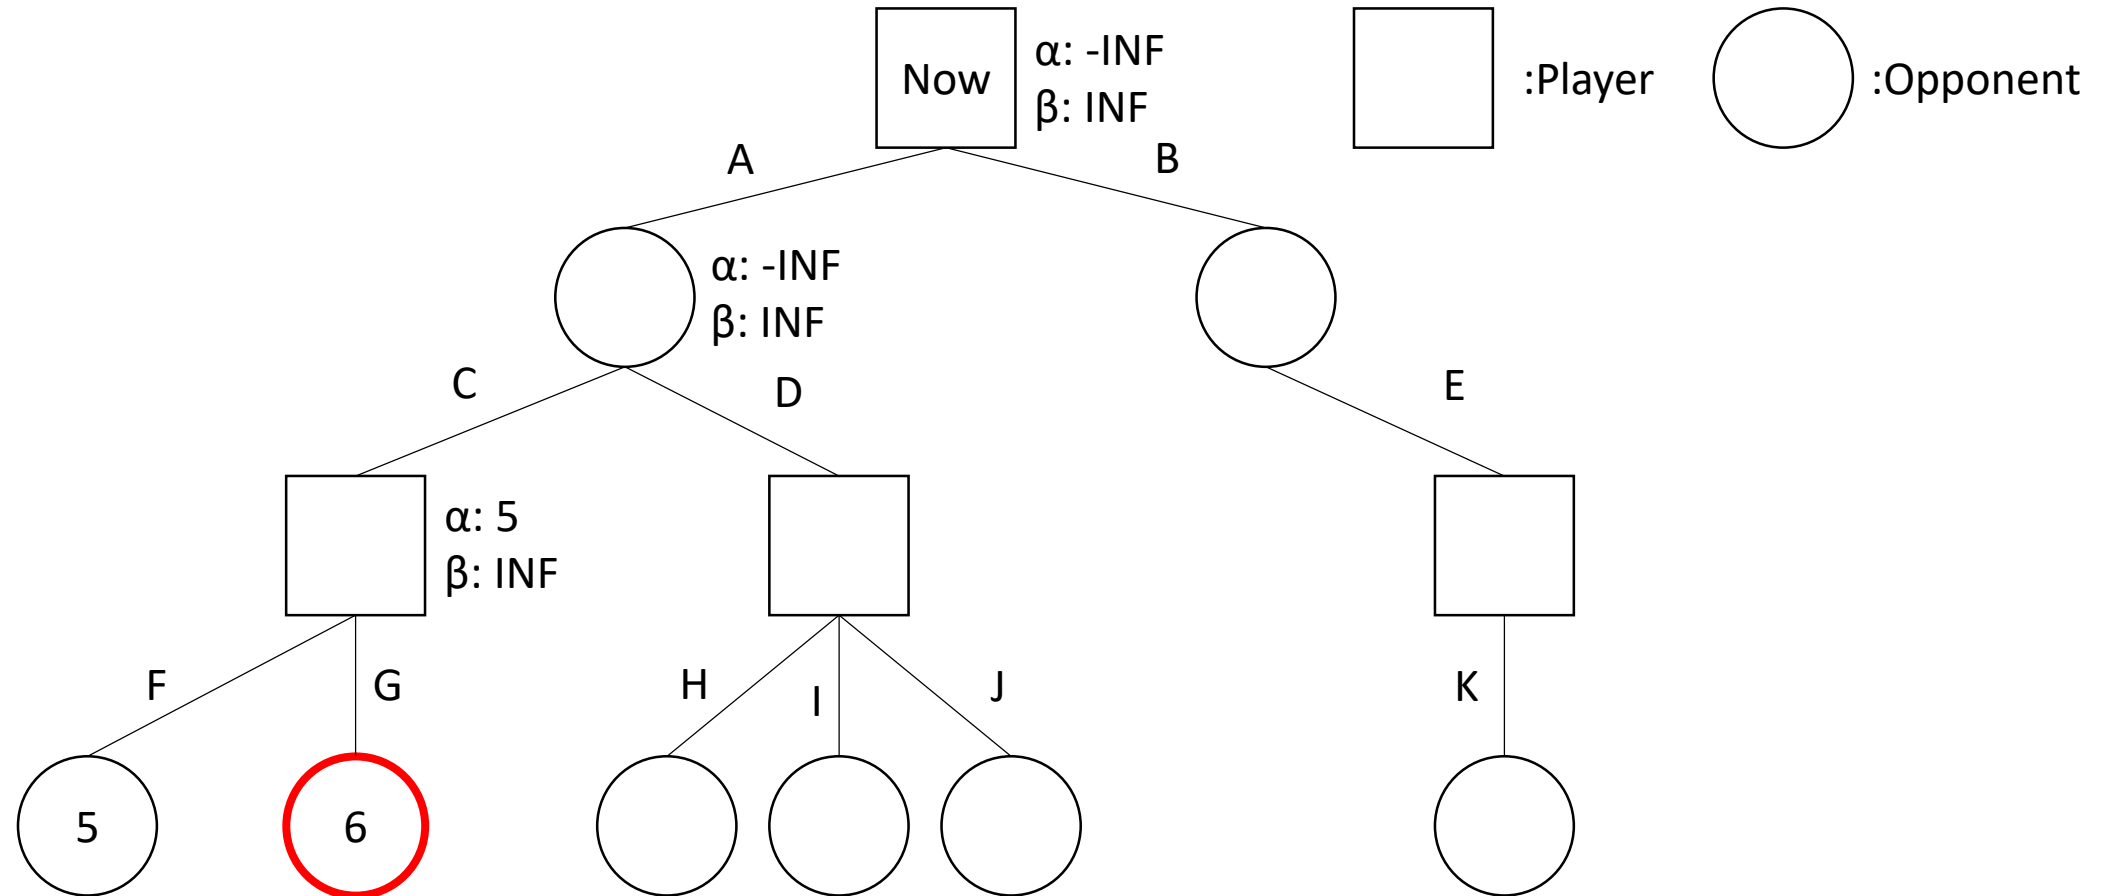
Evaluate score at leaves



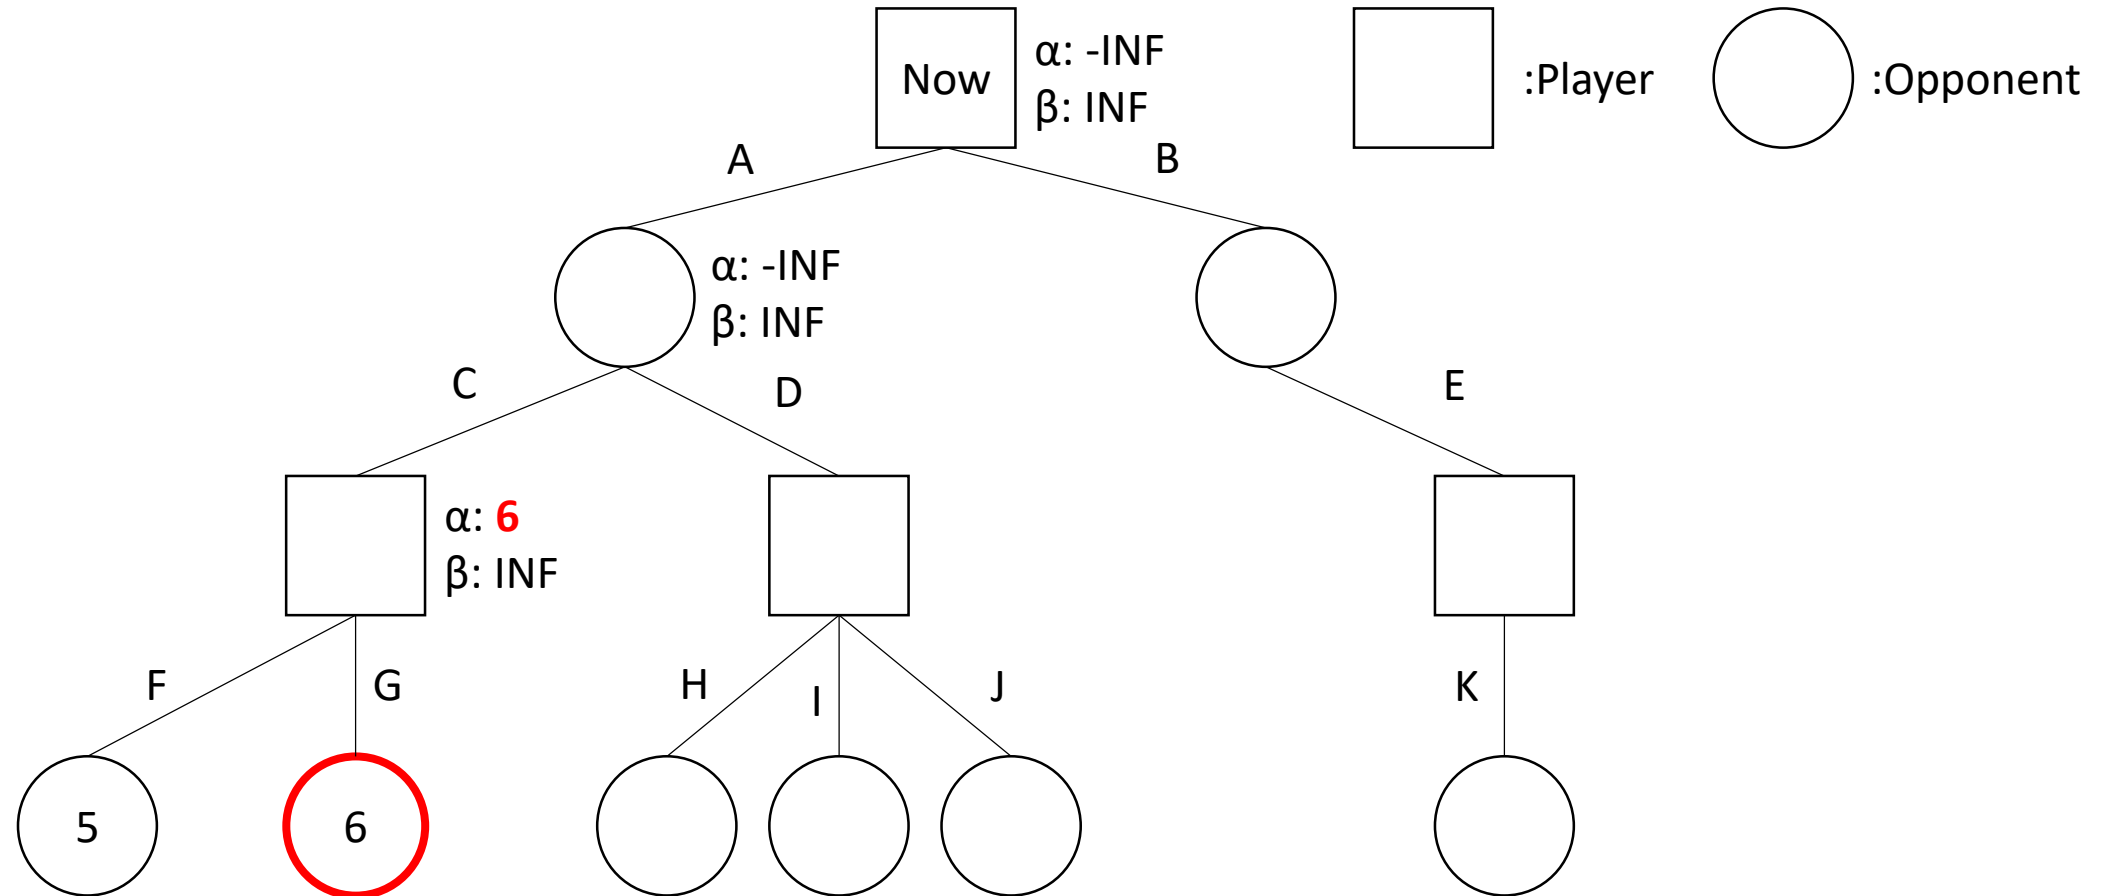
Update alpha



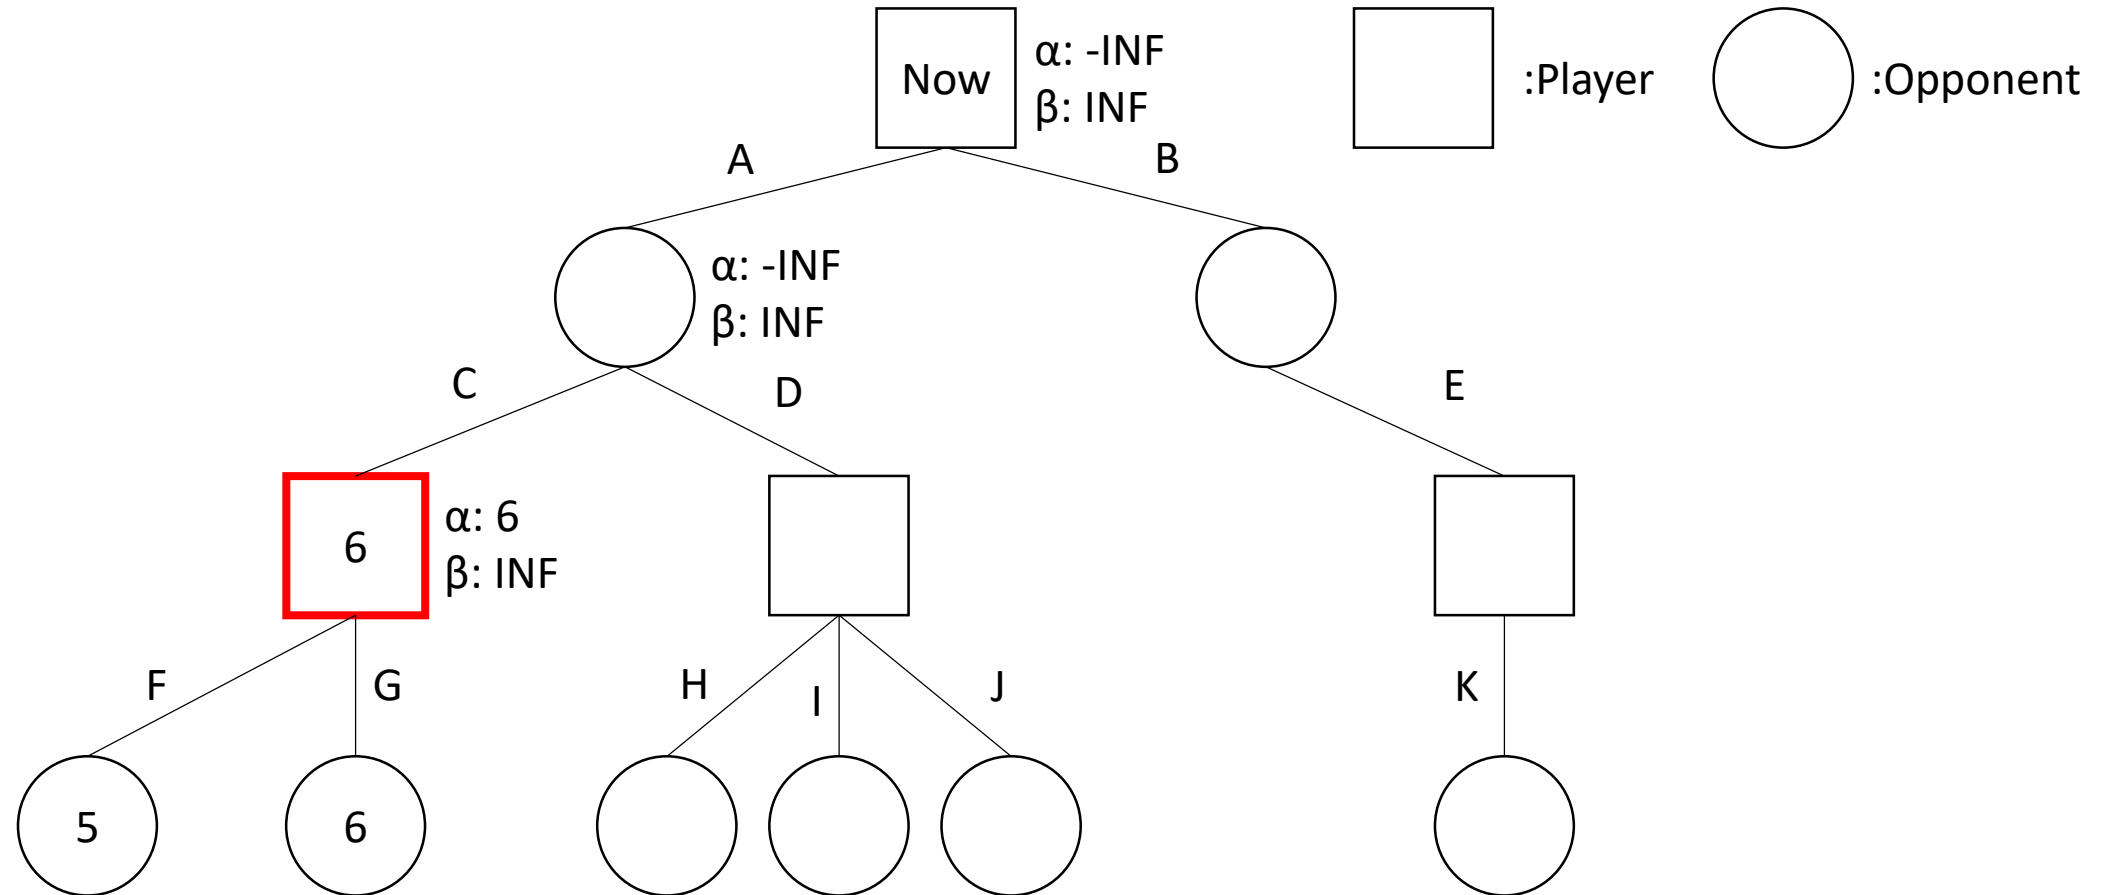
Evaluate score at leaves



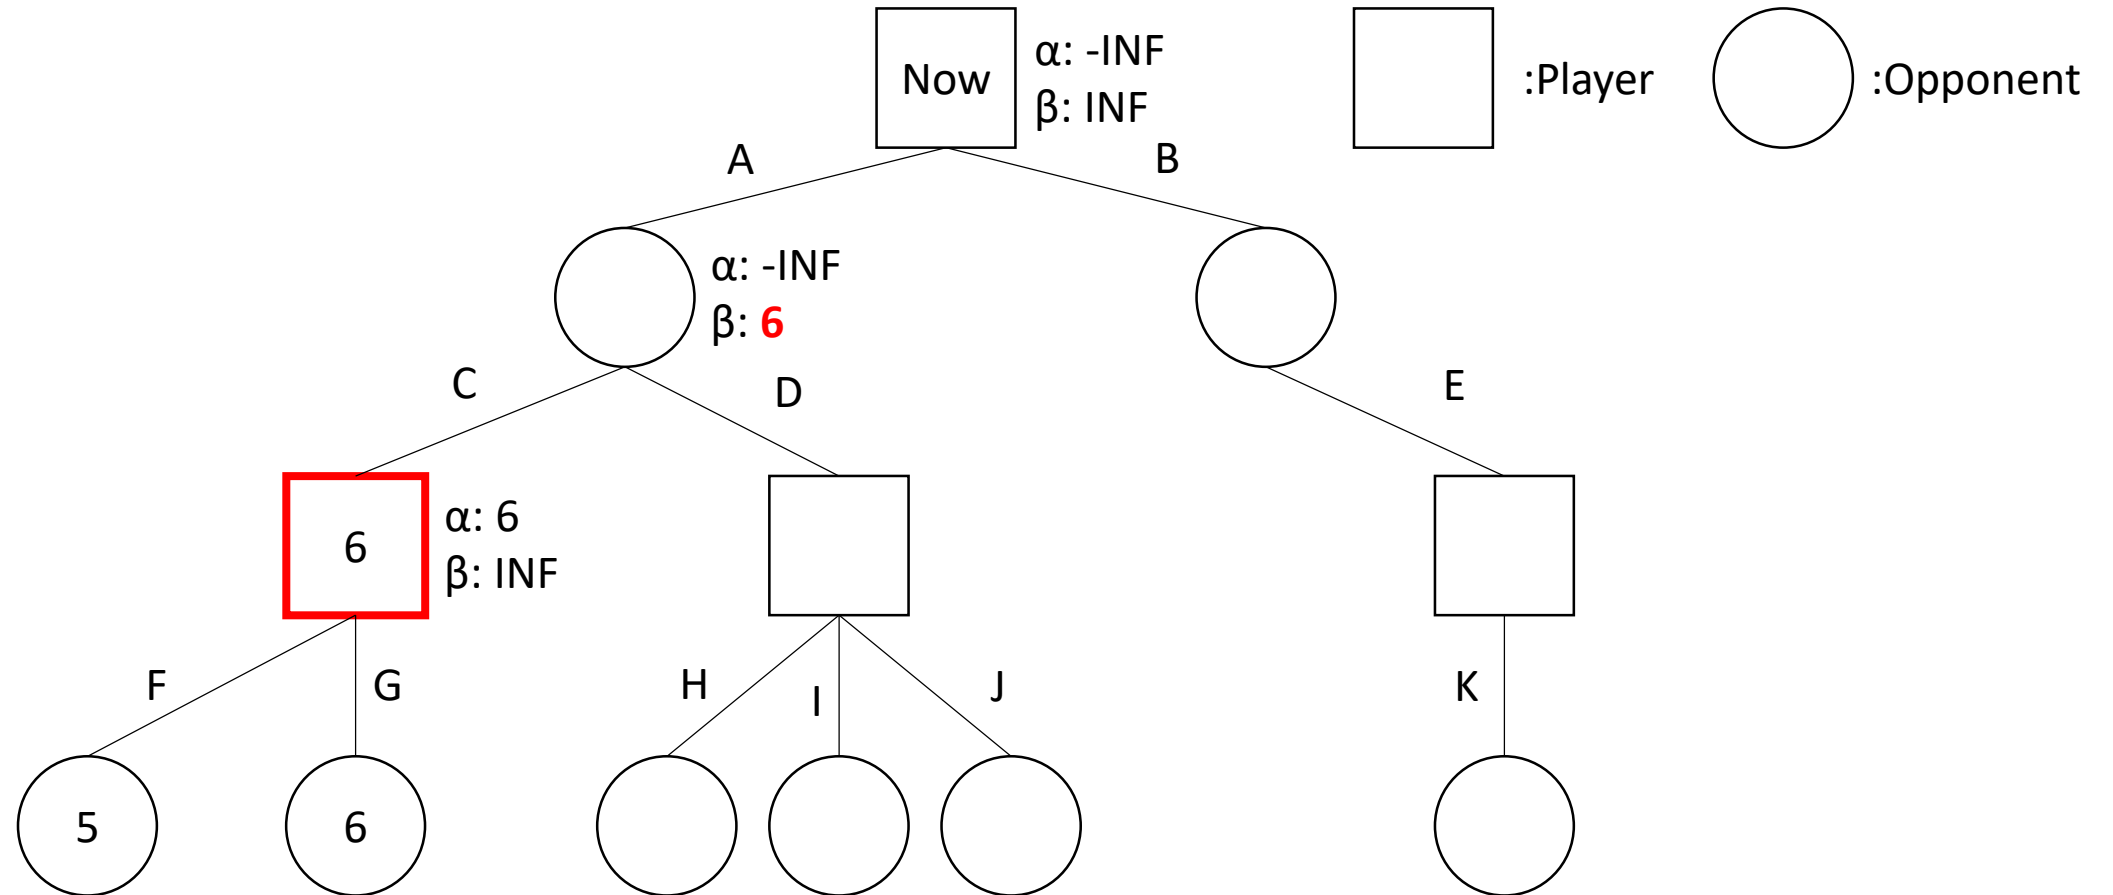
Update alpha



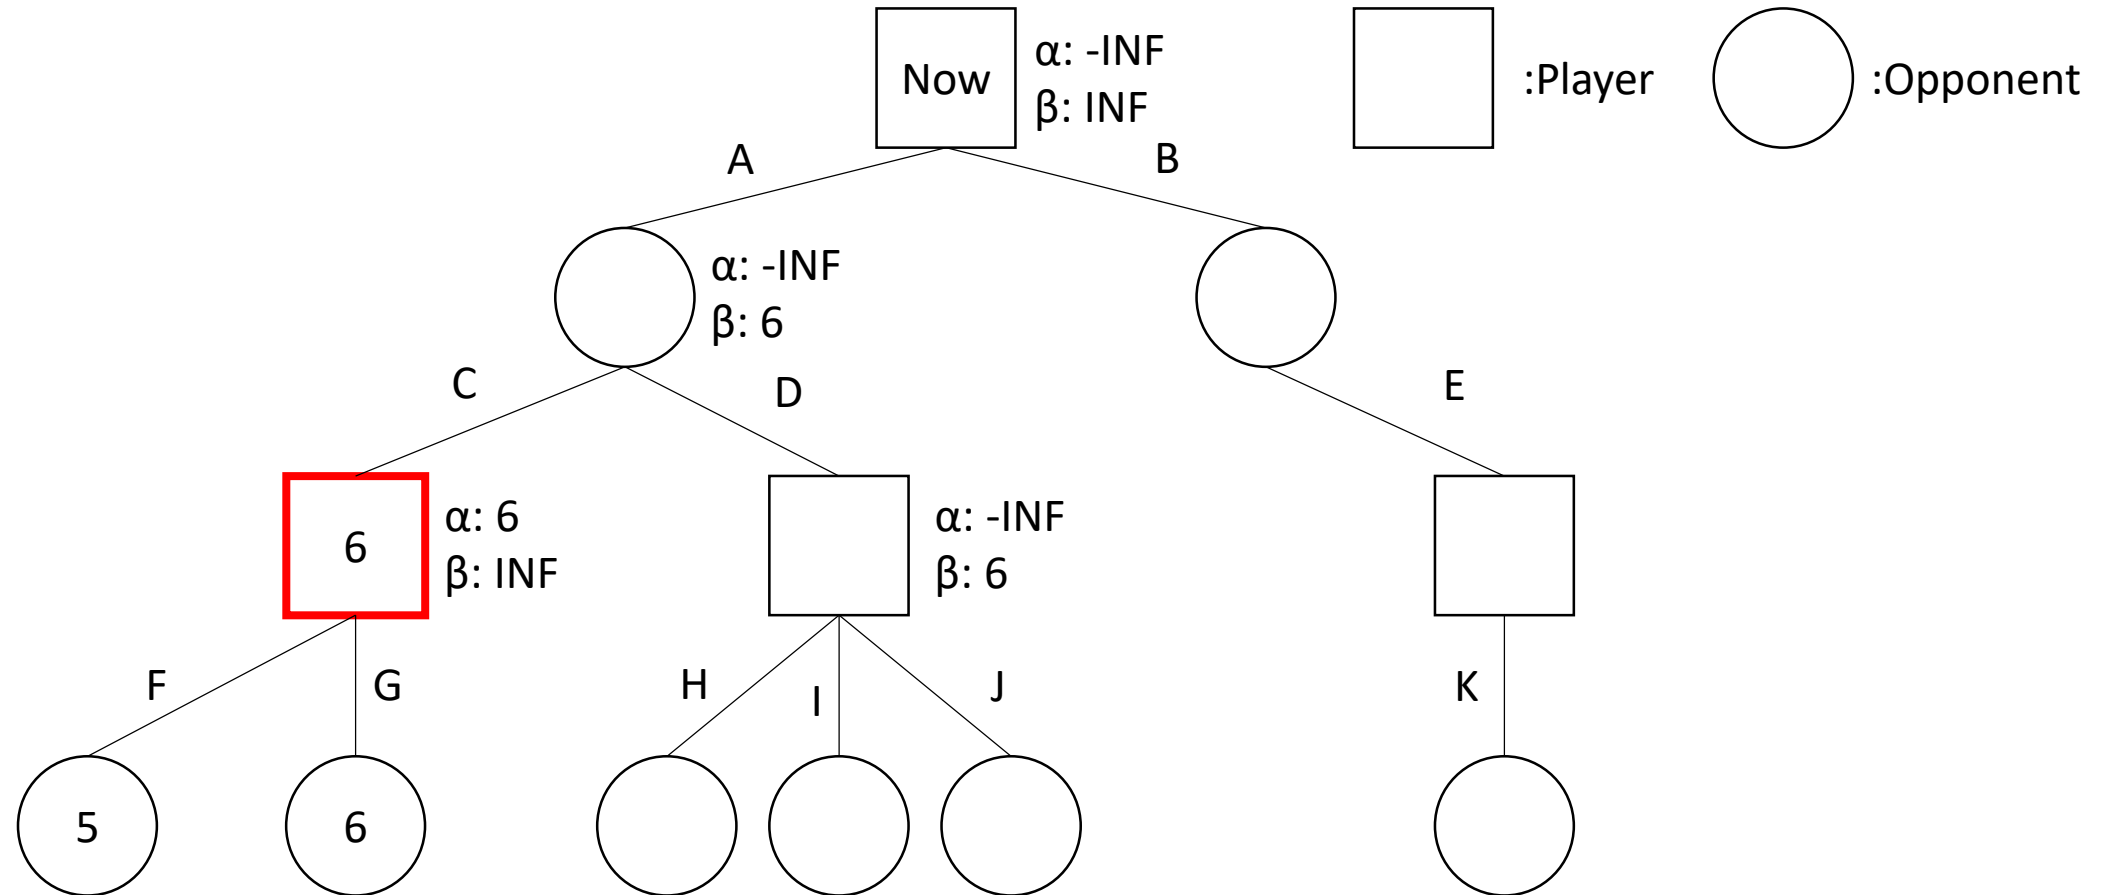
Player picks the largest score



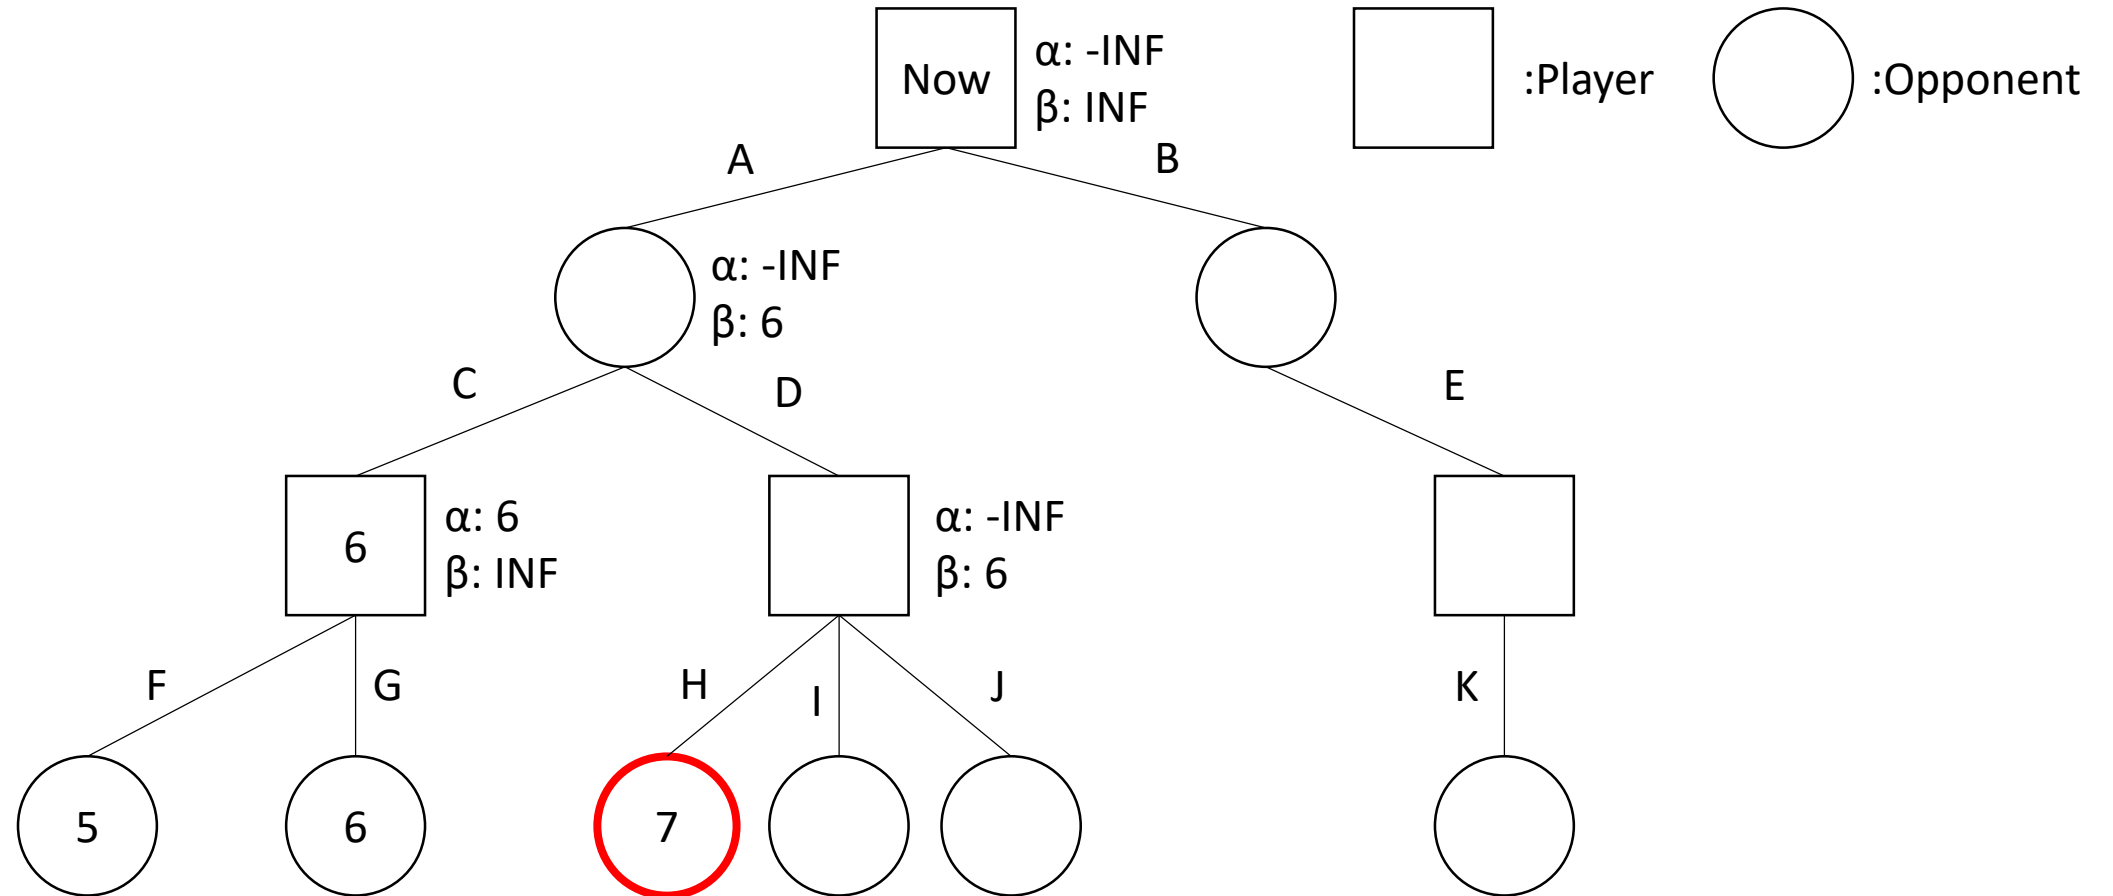
Update beta



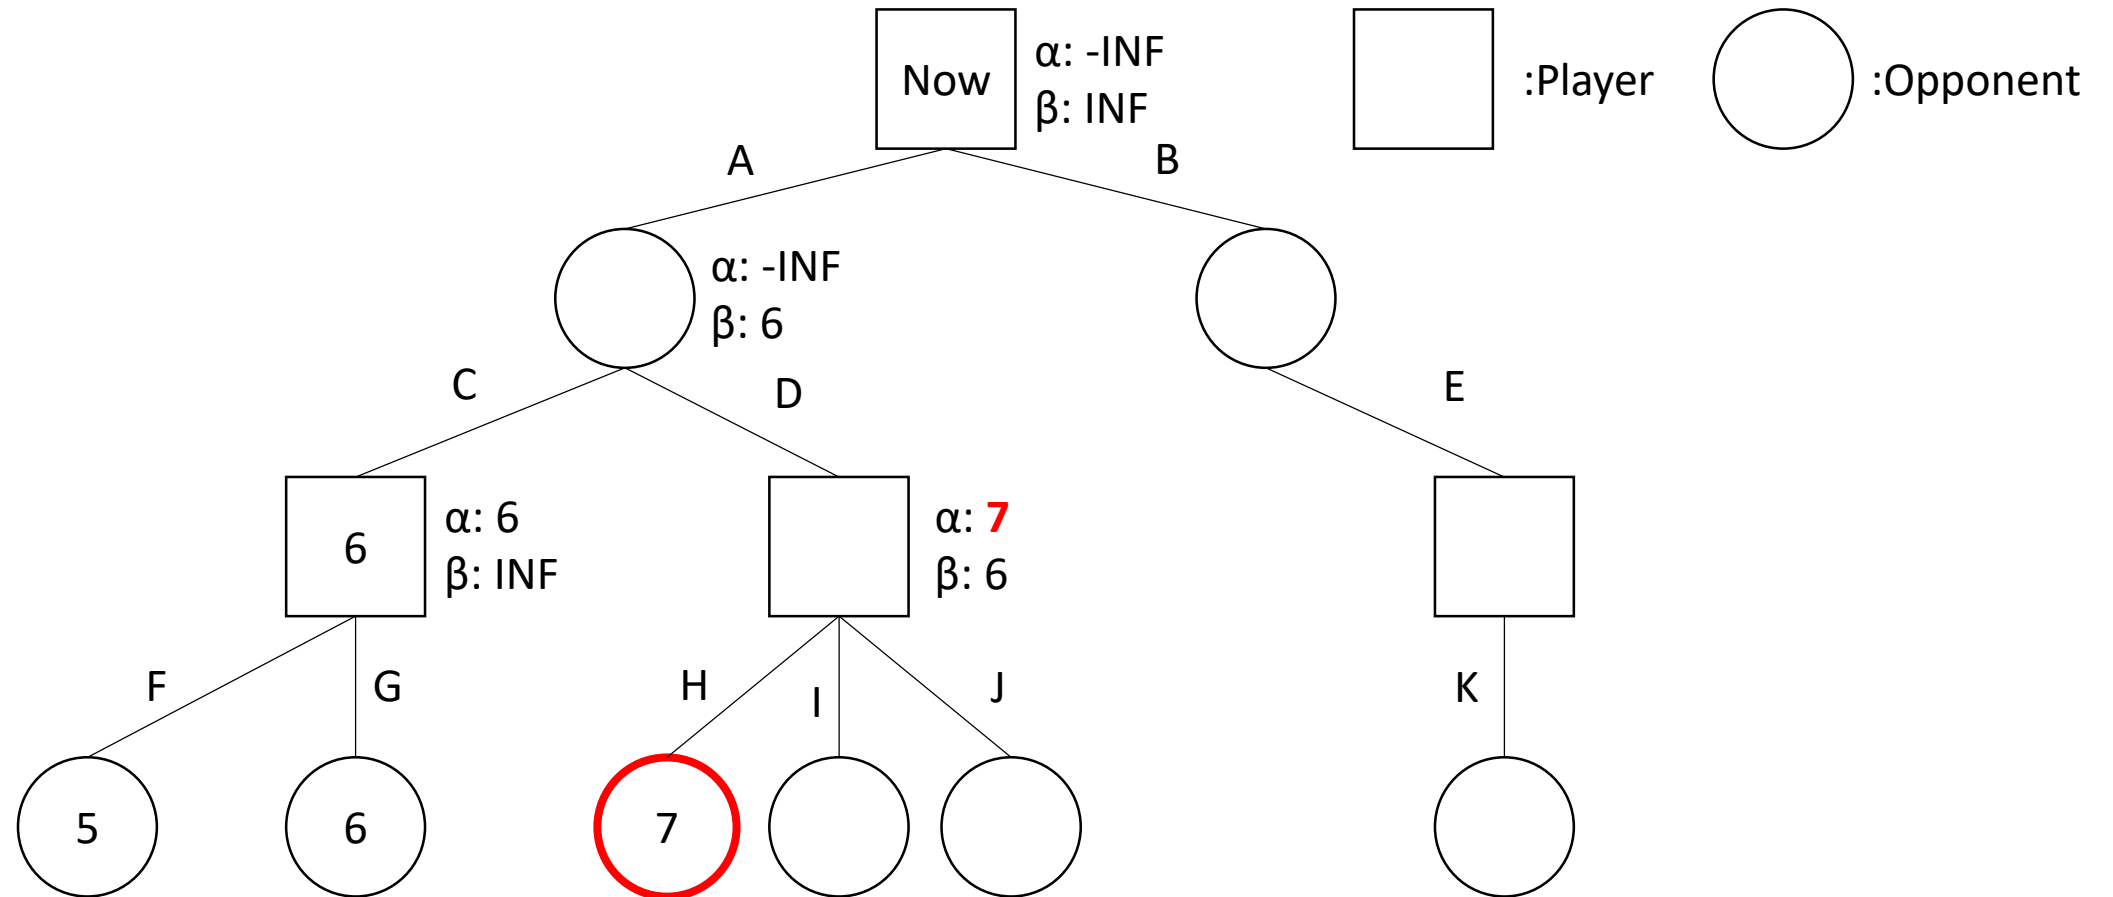
Propagate alpha and beta values



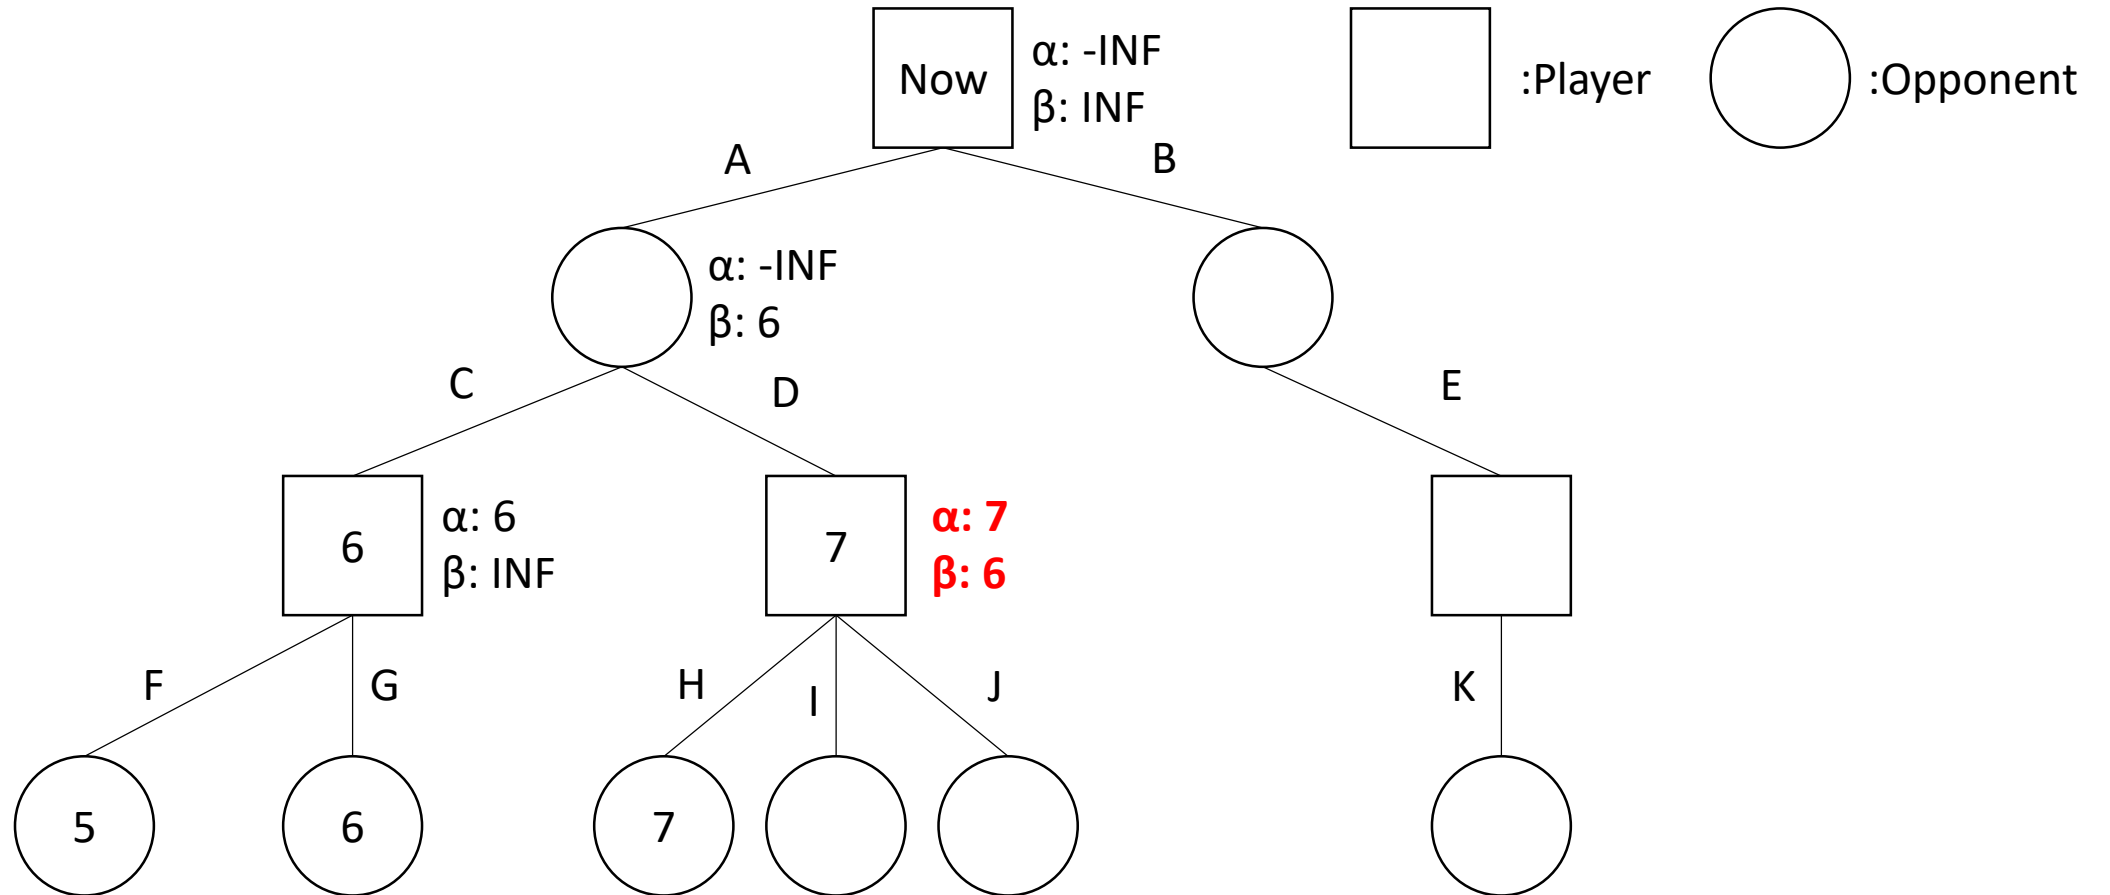
Evaluate score at leaves



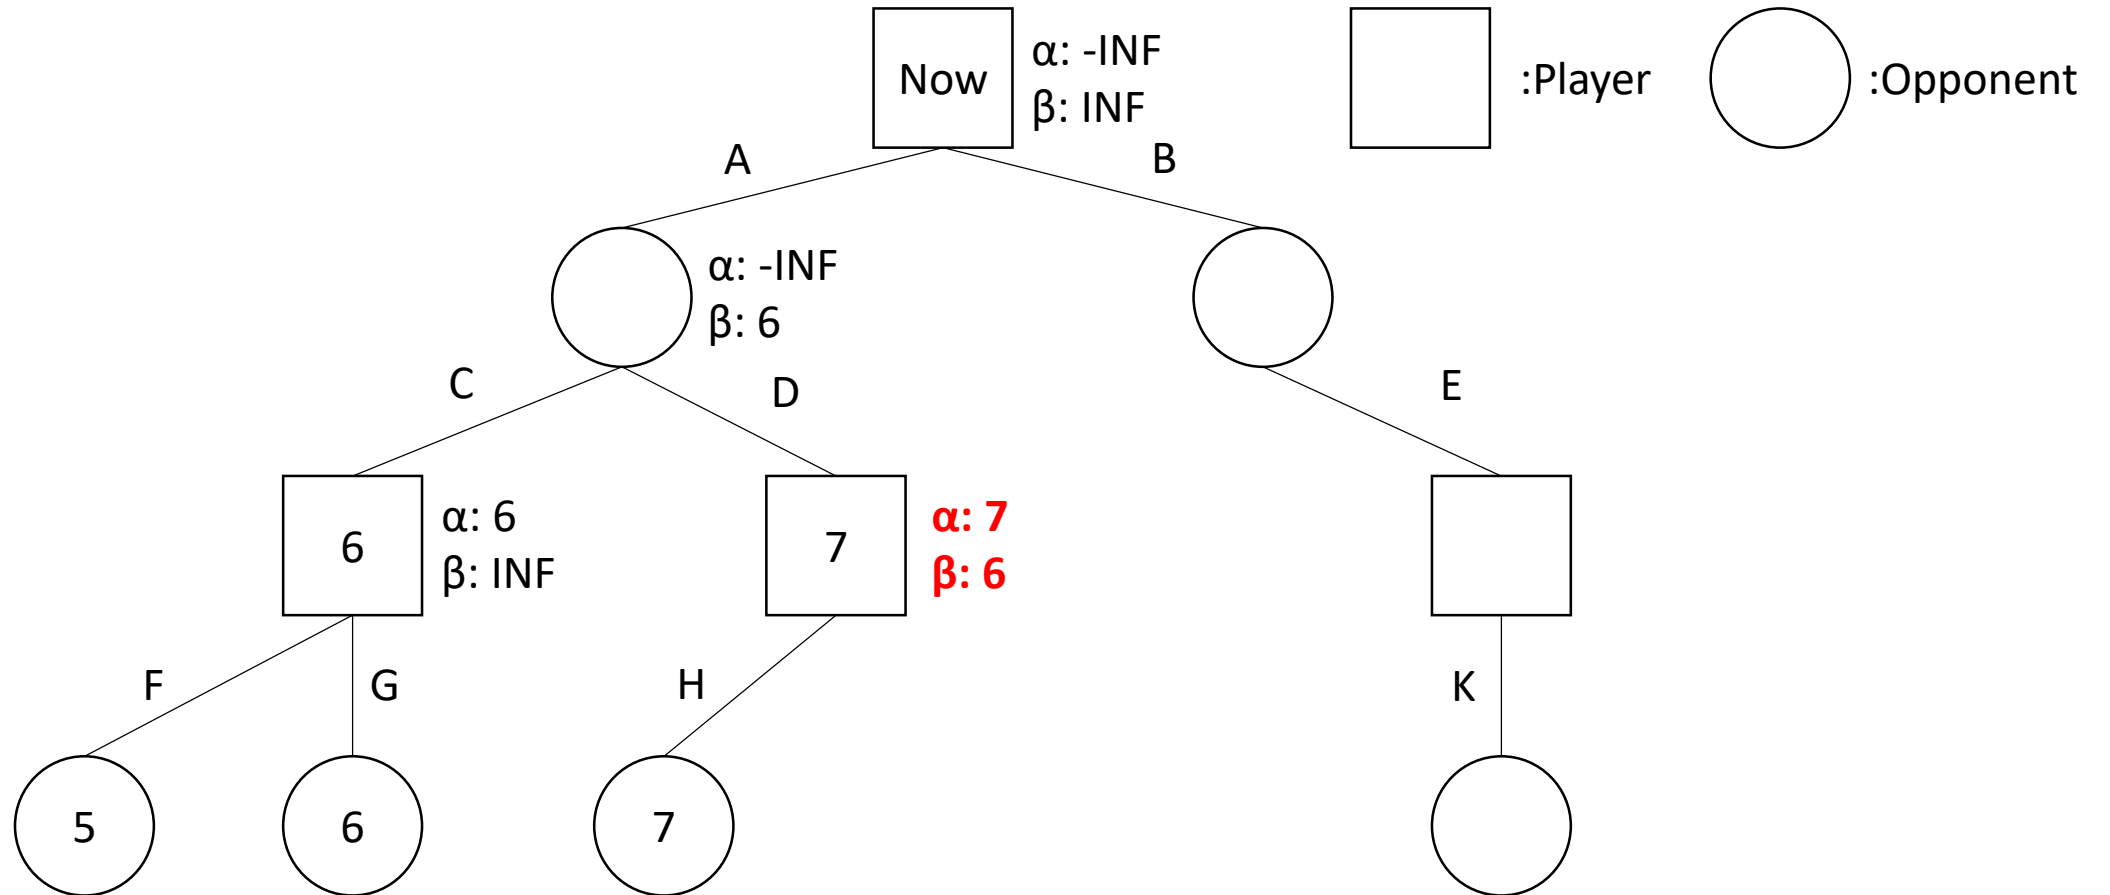
Update alpha



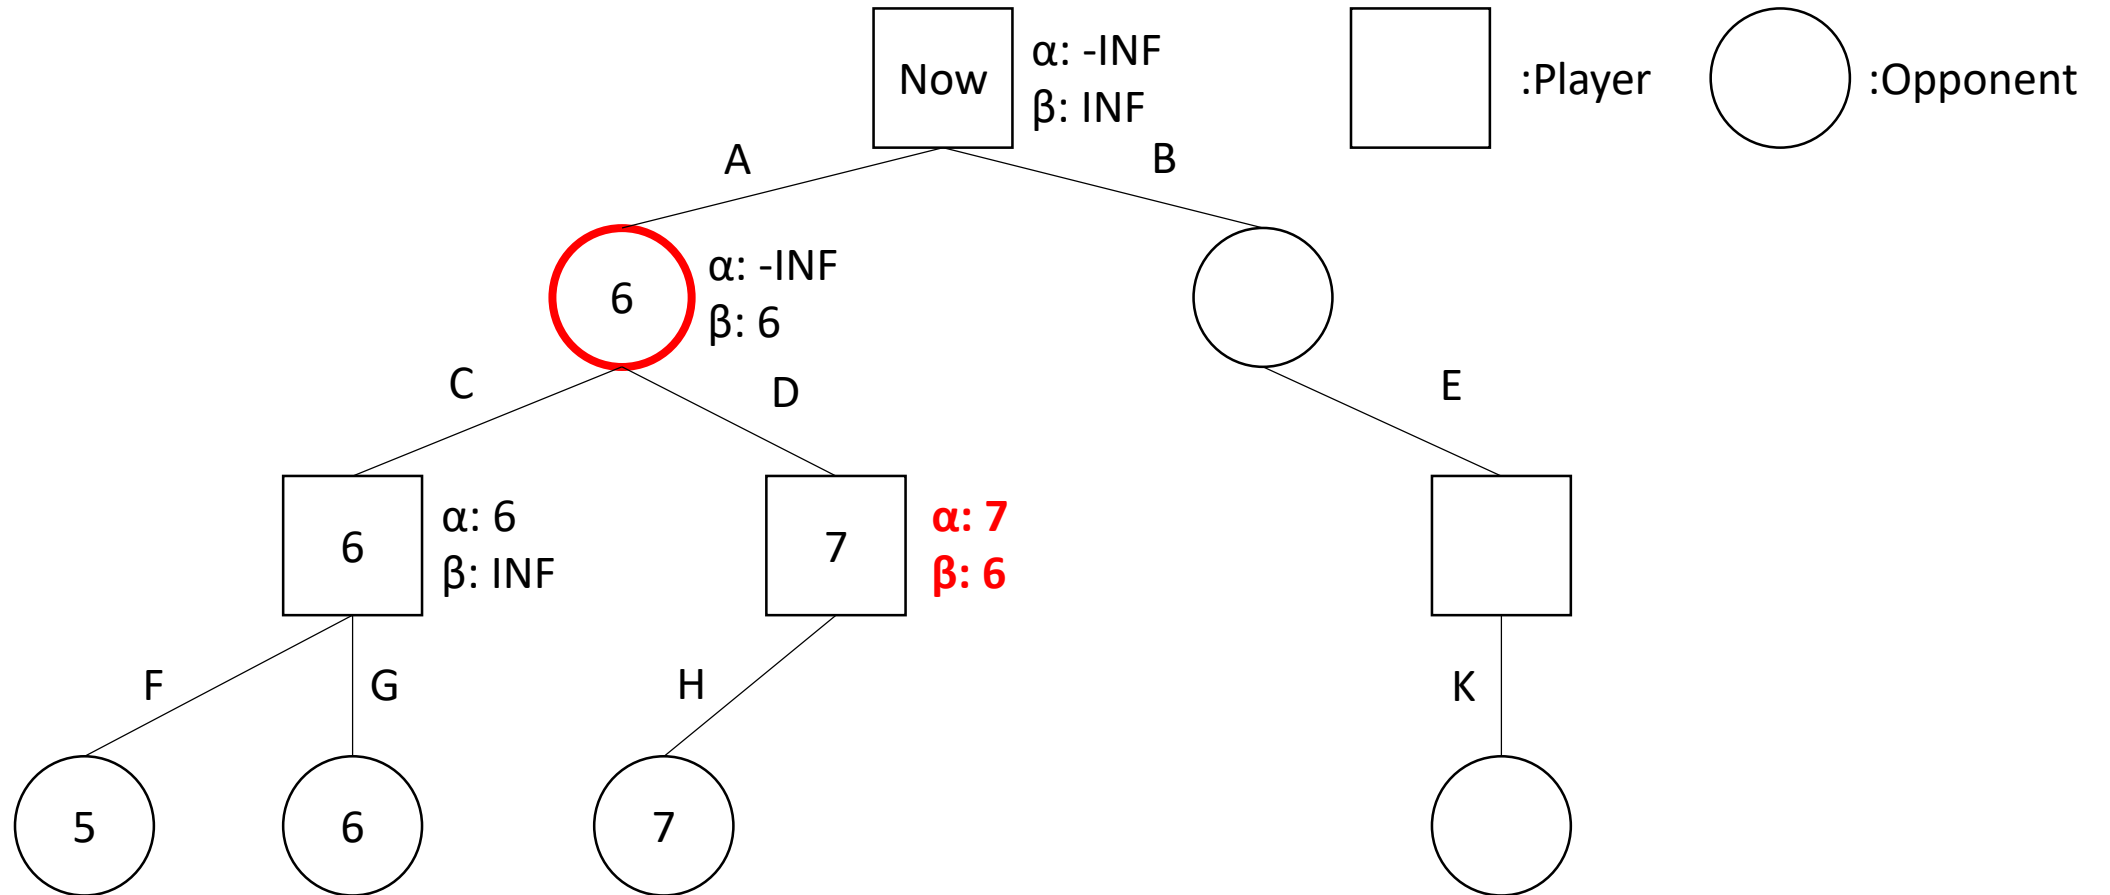
Alpha \geq beta in a player node, stop searching



Alpha \geq beta in a player node, stop searching



Opponent picks the smallest score



Alpha-Beta Pruning

- In the example above, we use the same search tree as Minimax
- By pruning, we eliminate branches I and J
- However, we still get the same result on branch A
- Alpha-Beta Pruning can effectively speed up the process while maintaining the same result

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
- 6. How To Design Your AI**
7. Package
8. Requirements
9. Grading

How To Design Your AI

- The game runner (main.cpp) executes the AIs of the player and the opponent in turns and communicates with them by files
- Your game AI should read the board status and available moves from the file “state”
- Your game AI should output your move to the file “action”

State file

- The state file consists of four parts:
- Current player (1 or 2)
- Board (8 x 8 matrix)
- Valid move number
- List of valid moves

```
1
0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0
0 0 0 2 1 0 0 0
0 0 0 1 2 1 0 0
0 0 0 0 0 2 0 0
0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0
4
2 3
3 2
5 4
6 5
```

Action file

- Your AI should output the next move to the “action” file
- You can keep output moves in the time limit (10 seconds)
- Only the last complete output will be considered
 - In the case on the right, (3, 5) will be accepted by the game runner
- You lose if you outputs an invalid move

```
4 2
7 1
3 5
3 5
2 2
2 4
2 4
3 5
3 5
```

How To Design Your AI

- You can refer to the “player_random.cpp” in the “src” folder of the package we provided
- Design your state value function to evaluate the board
- Implement the Alpha-Beta Pruning method and use your value function in the search process
- Run Alpha-Beta Pruning and decide which move to output

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
- 7. Package**
8. Requirements
9. Grading

Package Directory Structure

- Mini Project 3
 - baseline
 - windows
 - baseline<i>.exe, i := {1, 2, 3, 4, 5}
 - linux
 - baseline<i>, i := {1, 2, 3, 4, 5}
 - mac
 - baseline<i>, i := {1, 2, 3, 4, 5}
 - src
 - main.cpp
 - makefile
 - player_<description>.cpp, description := {random, infinite, partial, invalid}
- Mini Project 3 Introduction.pptx
- Mini Project 3 Spec.pdf

Baseline

- The baseline folder contains executables of the baselines
- If you use Windows, use the baselines in the windows folder
- If you use Linux, use the baselines in the linux folder
- If you use Mac, use the baselines in the mac folder
- We will use these baselines to test the performance of your program
- Test your program against these baselines and try to defeat them

Baseline

- baseline1 => pure random
- baseline2 => minimax with basic value function
- baseline3 => deeper minimax with basic value function
- baseline4 => alpha beta pruning with advanced value function
- baseline5 => deeper alpha beta pruning with advanced value function

Src

- The src folder contains the main program, a makefile and some provided example programs
- The main.cpp is the game runner, **do not modify it**
 - However, you can refer to it when implementing game simulation of minimax
- The makefile helps you compile your code
 - You can also compile by hand, **make sure you set the flags -Wall and -Wextra**

Src

- The example programs are some situations you might face while testing your program
- `player_random.cpp` => the pure random AI
 - You can refer to this one and write your own code
- `player_infinite.cpp` => only the last move you output in your turn will be considered
- `player_invalid.cpp` => you lose if you output an invalid move
- `player_partial.cpp` => moves in incomplete format will not be considered

How To Compile Your Code

- To compile all programs:
- Type **make** when your working directory is 'src'

- To compile <program_name>.cpp:
- Type **make <program_name>** when your working directory is 'src'

- To clean the executables:
- Type **make clean** when your working directory is 'src'

How To Test Your Code

- If you use Windows, type the command below in cmd:
 - `main.exe <AI1>.exe <AI2>.exe`
- If you use Linux / Max, type the command below in terminal:
 - `./main ./<AI1> ./<AI2>`

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading

Requirements - Code

- Design and implement an AI that can play the boardgame Othello
- Your AI should read the board from file and write next move to file
- Design a value function to evaluate the score of the board
- Enhance the policy of your AI with Alpha-Beta Pruning

Requirements - Code

- If you are not satisfied by the Alpha-Beta Pruning algorithm, you can try some more advanced methods. However, make sure you can explain how Minimax and Alpha-Beta Pruning works during demo
- If you cannot complete the Alpha-Beta Pruning algorithm, implementing the basic Minimax algorithm also gives you some score

Requirements - Code

- You will lose immediately if your program outputs an invalid move
- Time limit for each move is 10 seconds, and the memory limit is 4GB
- You can keep output moves in the limited time. Only the last successful output move is used by the game runner
- Please refer to the spec for more detailed rules

Requirements - Submission

- Please use **C++** and write your program in a **single file**
- Your program should be named as **<student_id>_project3.cpp**
- Your program will be compiled in a GNU / Linux environment by:
 - **g++ -O2 -std=c++14 -Wall -Wextra <student_id>_project3.cpp**
- Please make sure your program can be compiled by the command above with no error

Requirements - Report and Demo

- You should write a report to elaborate on how you design your AI
- The report is not directly graded, but is your **only available reference through the TA demo** (You cannot refer to your code in demo)
- You must attend the demo and answer the questions from TA
- **The demo date and method will be announced soon**

Outline

1. Introduction
2. Othello
3. State Value Function
4. Minimax
5. Alpha-Beta Pruning
6. How To Design Your AI
7. Package
8. Requirements
9. Grading


Grading

- The project accounts for 9 points of your total grade
- Beat every baseline => +5 points (1 point for each baseline)
- Implement Tree search (Minimax) => +2 points
- Design of your state value function => +1 point
- Implement Alpha-Beta Pruning => +1 point

Grading Bonus

- (Bonus) Uses version control software => +1 point
- Include a screenshot with more than 3 commits in your report
- (Bonus) Class ranking => At most +3 points
- You can attend the class ranking if you beat all baselines
- Your AI will play against other AIs of your classmates and gain bonus score according to your ranking

Develop four own AI!



AlphaOthello